



中山大學  
SUN YAT-SEN UNIVERSITY

数据科学与计算机学院  
School of Data and Computer Science

# 嵌入式系统导论

习题3：资源共享



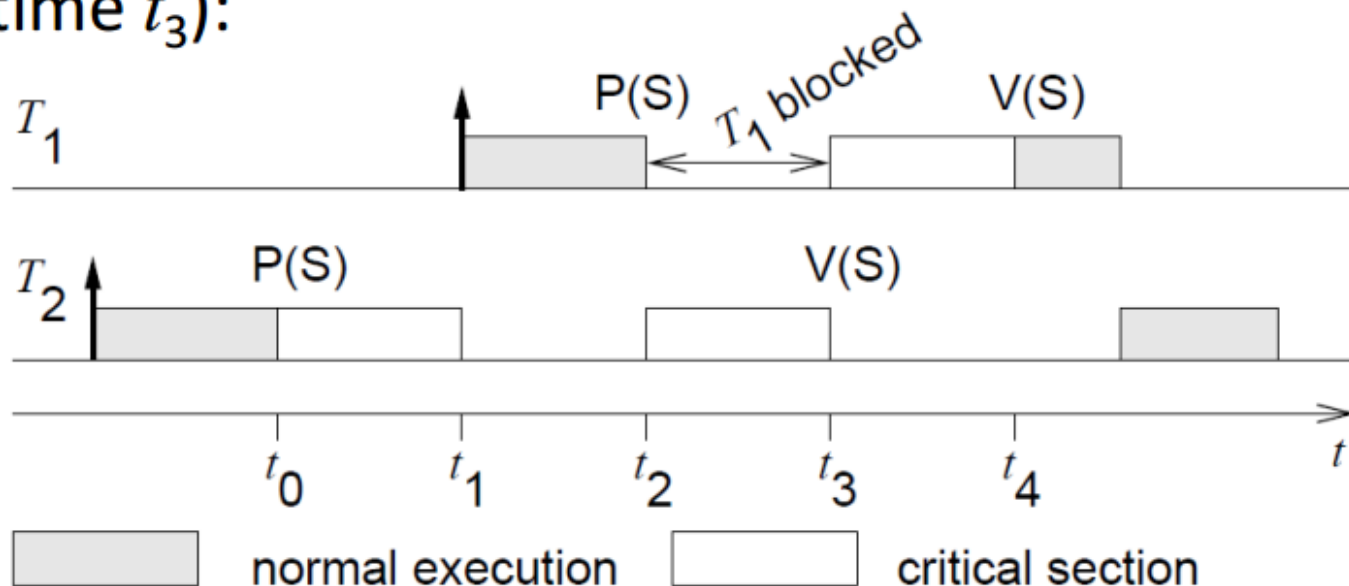
# Terms

- Each **exclusive resource**  $R_i$  must be protected by a different **semaphore**  $S_i$  and each critical section operating on a resource must begin with a  $\text{wait}(S_i)$  primitive and end with a  $\text{signal}(S_i)$  primitive.
- All tasks blocked on the same resource are kept in a queue associated with the semaphore. When a running task executes a **wait** on a **locked semaphore**, it enters a **waiting state**, until another task executes a **signal** primitive that **unlocks the semaphore**.



# Priority Inversion ( 1 )

- Priority  $T_1$  assumed to be  $>$  than priority of  $T_2$ .
- If  $T_2$  requests exclusive access first (at  $t_0$ ),  $T_1$  has to wait until  $T_2$  releases the resource (at time  $t_3$ ):

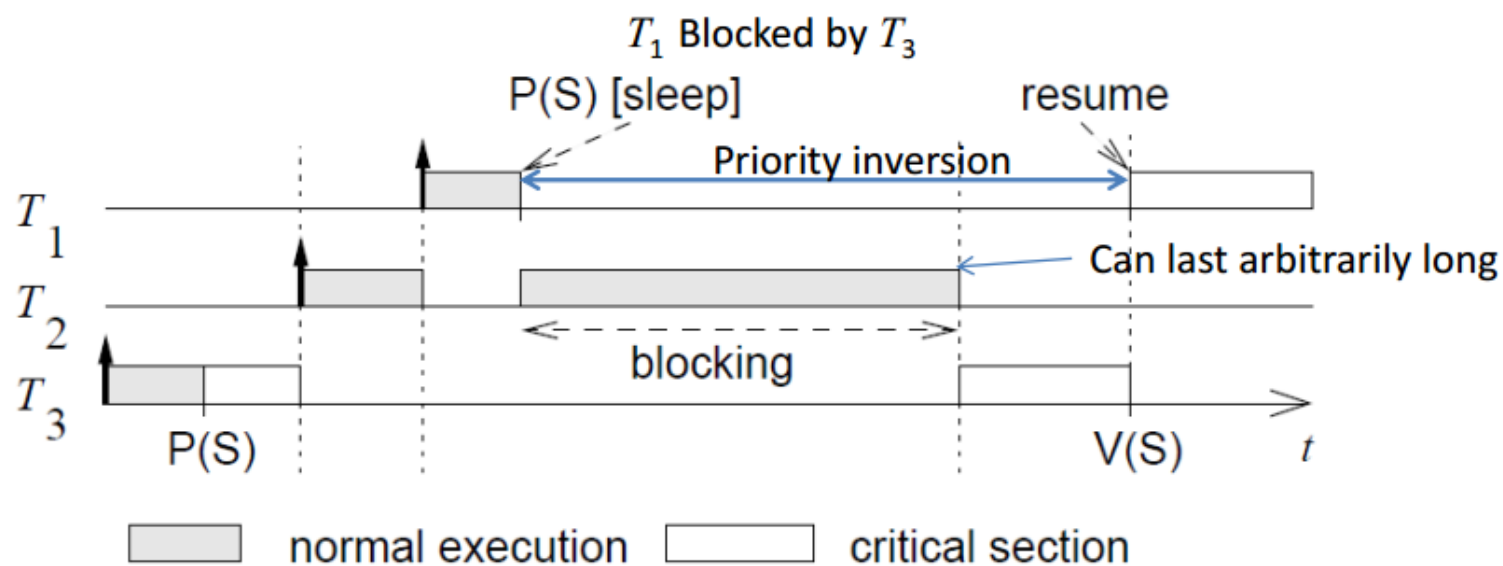


For 2 tasks: blocking is bounded by the length of the critical section



# Priority Inversion ( 2 )

- Blocking with  $>2$  tasks can exceed the length of any critical section
  - Priority of  $T_1 >$  priority of  $T_2 >$  priority of  $T_3$ .
  - $T_2$  preempts  $T_3$ :  $T_2$  can prevent  $T_3$  from releasing the resource.





# Resource Access Protocols

- Basic idea: Modify the priority of those tasks that cause blocking. When a task  $T_i$  blocks one or more higher priority tasks, it temporarily assumes a higher priority.
- Methods:
  - Priority Inheritance Protocol (PIP), for static priorities
  - Priority Ceiling Protocol (PCP), for static priorities
  - Stack Resource Policy (SRP),
    - For static and dynamic priorities
  - others ...



# Priority Inheritance Protocol(PIP)

- Assumptions:
  - $n$  tasks which cooperate through  $m$  shared resources; fixed priorities, all critical sections on a resource begin with a `wait( $S_i$ )` and end with a `signal( $S_i$ )` operation.
- Basic idea:
  - When a task  $T_i$  blocks one or more higher priority tasks, it temporarily assumes (inherits) the **highest** priority of the blocked tasks.
- Terms:
  - We distinguish a fixed **nominal priority**  $P_i$  and an **active priority**  $p_i$  larger or equal to  $P_i$ . Tasks  $T_1, \dots, T_n$  are ordered with respect to nominal priority where  $T_1$  has **highest priority**. Tasks do not suspend themselves.





# PIP: Algorithm


- Tasks are scheduled according to their **active priorities**. Tasks with the same priorities are scheduled FCFS.
- When a task  $T_i$  tries to **enter a critical section** and the resource is blocked by a lower priority task, the task  $T_i$  is blocked. Otherwise it enters the critical section.
- When a task  $T_i$  is **blocked**, it transmits its active priority to the task  $T_k$  that holds the semaphore.  $T_k$  resumes and executes the rest of its critical section with a priority  $p_k = p_i$  (it **inherits** the priority of the highest priority of the tasks blocked by it).
- When  $T_k$  exits a critical section, it **unlocks** the semaphore and the highest priority task blocked on that semaphore is awakened. If no other tasks are blocked by  $T_k$ , then  $p_k$  is set to  $P_k$ , otherwise it is set to the highest priority of the tasks blocked by  $T_k$ .
- Priority inheritance is **transitive**, i.e., if 1 is blocked by 2 and 2 is blocked by 3, then 3 inherits the priority of 1 via 2.




# Example:

- How would priority inheritance affect our example with 3 tasks?

Before:



PIP:



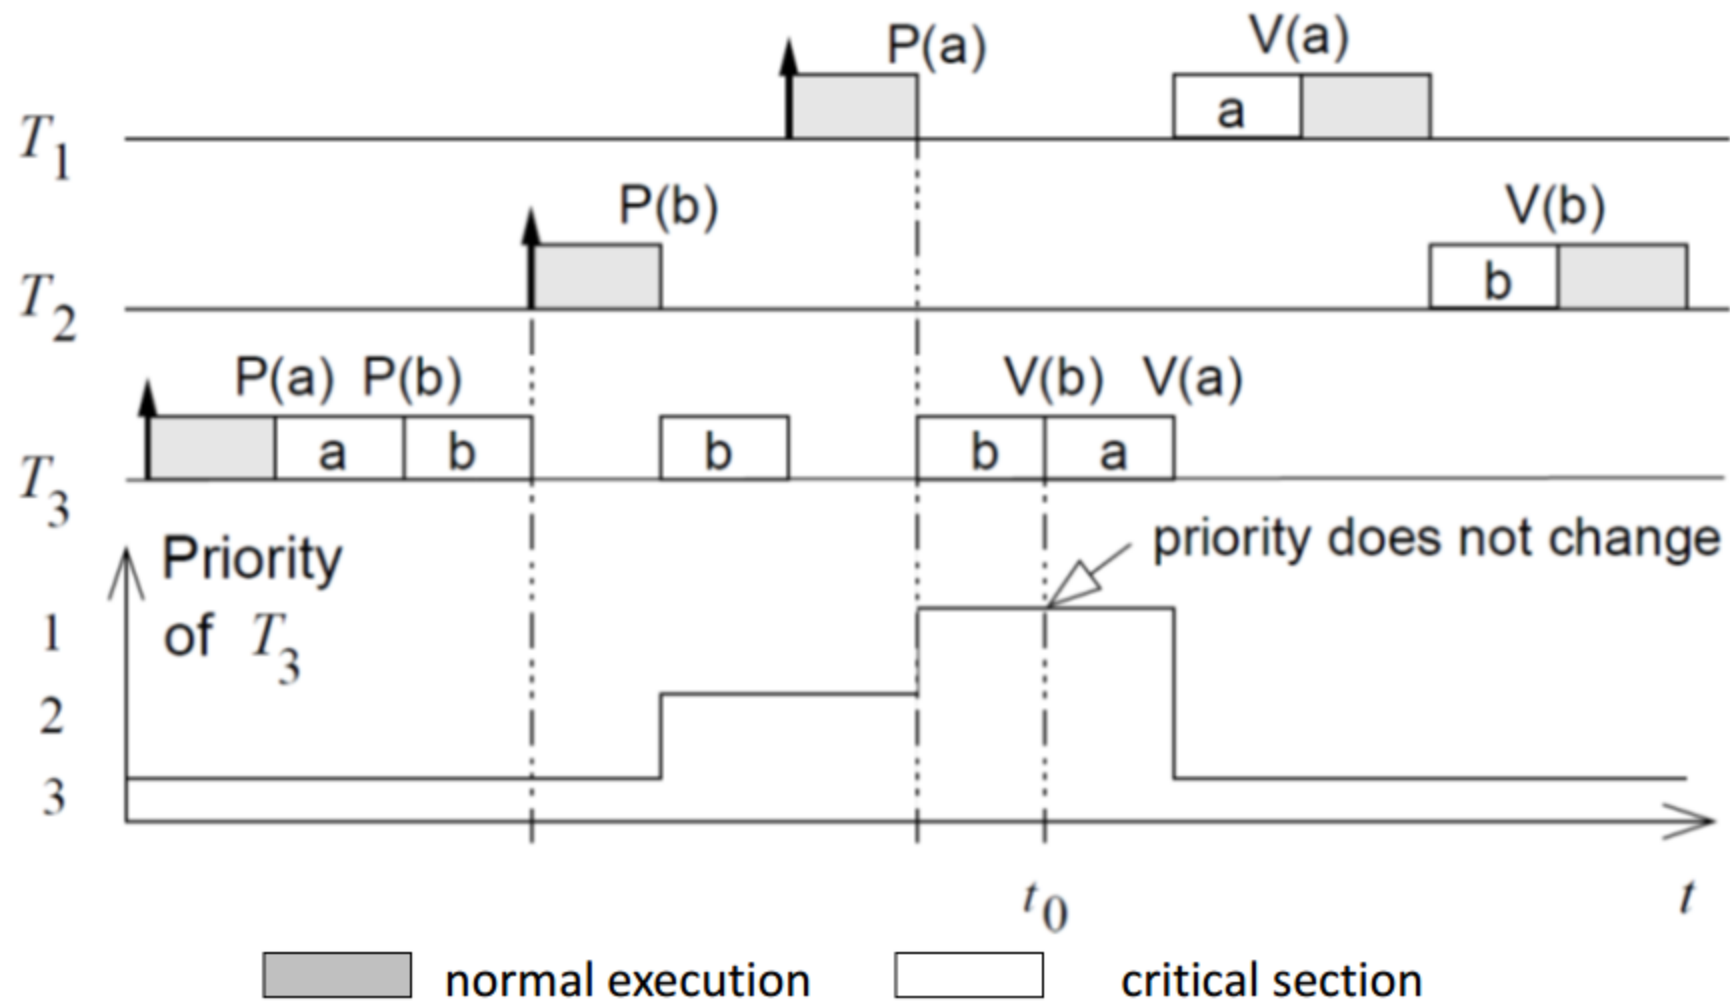
$T_3$  inherits the priority of  $T_1$  and  $T_3$  resumes.

- Direct Blocking:** higher-priority task tries to acquire a resource held by a lower-priority task
- Push-through Blocking:** medium-priority task is blocked by a lower-priority. Task that has inherited a higher priority from a task it directly blocks



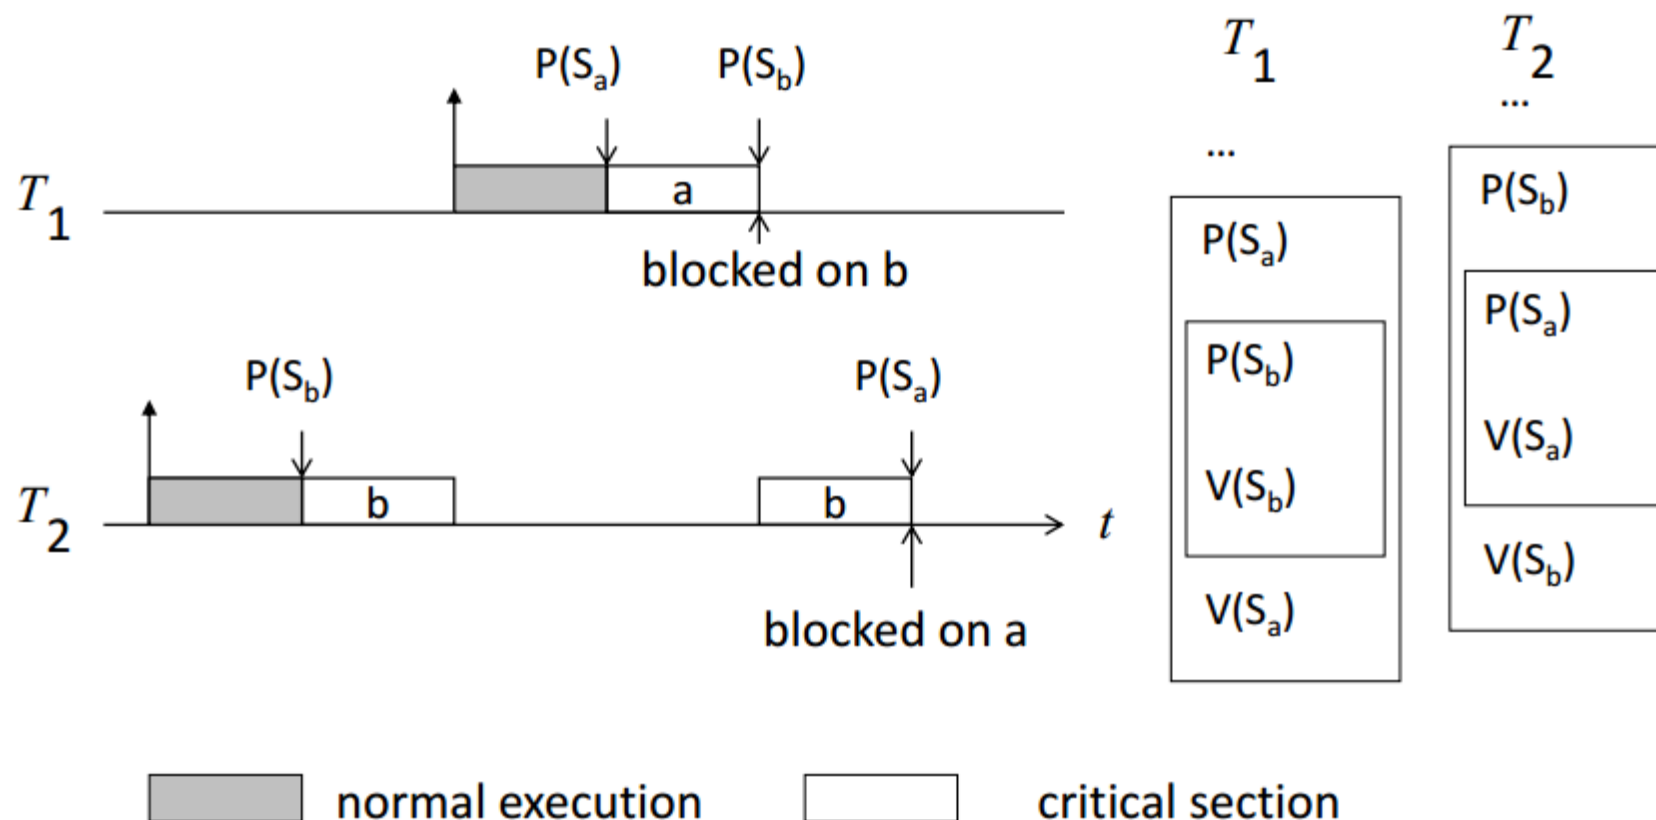


# Example:





# Deadlock is Possible



- Problem exists also when no priority inheritance is used



# PIP--->PCP

## ■ The Priority Inheritance Protocol (PIP)

- does not prevent deadlocks
- can lead to chained blocking
  - (Several lower priority tasks can block a higher priority task)
- and has inherent static priorities of tasks

## → The Priority Ceiling Protocol (PCP)

- avoids multiple blocking
- guarantees that, once a task has entered a critical section, it cannot be blocked by lower priority tasks until its completion.



# PCP


- A task is not allowed to enter a critical section if there are already locked semaphores which could block it eventually
- Hence, once a task enters a critical section, it can not be blocked by lower priority tasks until its completion.
- This is achieved by assigning priority ceiling.
- Each semaphore  $S_k$  is assigned a **priority ceiling**  $C(S_k)$ . It is the priority of the highest priority task that can lock  $S_k$ .

This is a static value.

A job  $j$  executing  $P(S)$  is granted access to  $S$  if the priority of  $j$  is strictly higher than the ceiling of any semaphore locked by a job other than  $j$ . Otherwise,  $j$  becomes blocked and  $S$  is not allocated to  $j$ .



# Example :



$C(S0=P1) \ C(S1=P1) \ C(S2=P2)$

Source: [http://www.ida.liu.se/unmbo/RTS\\_CUGS\\_files/Lecture3.pdf](http://www.ida.liu.se/unmbo/RTS_CUGS_files/Lecture3.pdf)



# 作业

- 11月30日晚12:00之前提交EX1到ftp相应文件夹