

Chapter 3: Processes



Chapter 3: Processes

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Cooperating Processes
- ❑ Interprocess Communication
- ❑ Communication in Client-Server Systems




Process Concept

- ❑ An operating system executes a variety of programs:
 - Batch system – jobs
 - Time-shared systems – user programs or tasks
- ❑ Textbook uses the terms **job** (作业) and **process** almost interchangeably
- ❑ **Process** – a program in **execution**; process execution must progress **in sequential fashion**
- ❑ A process includes:
 - **p**rogram **c**ounter
 - stack
 - data section



Process in Memory

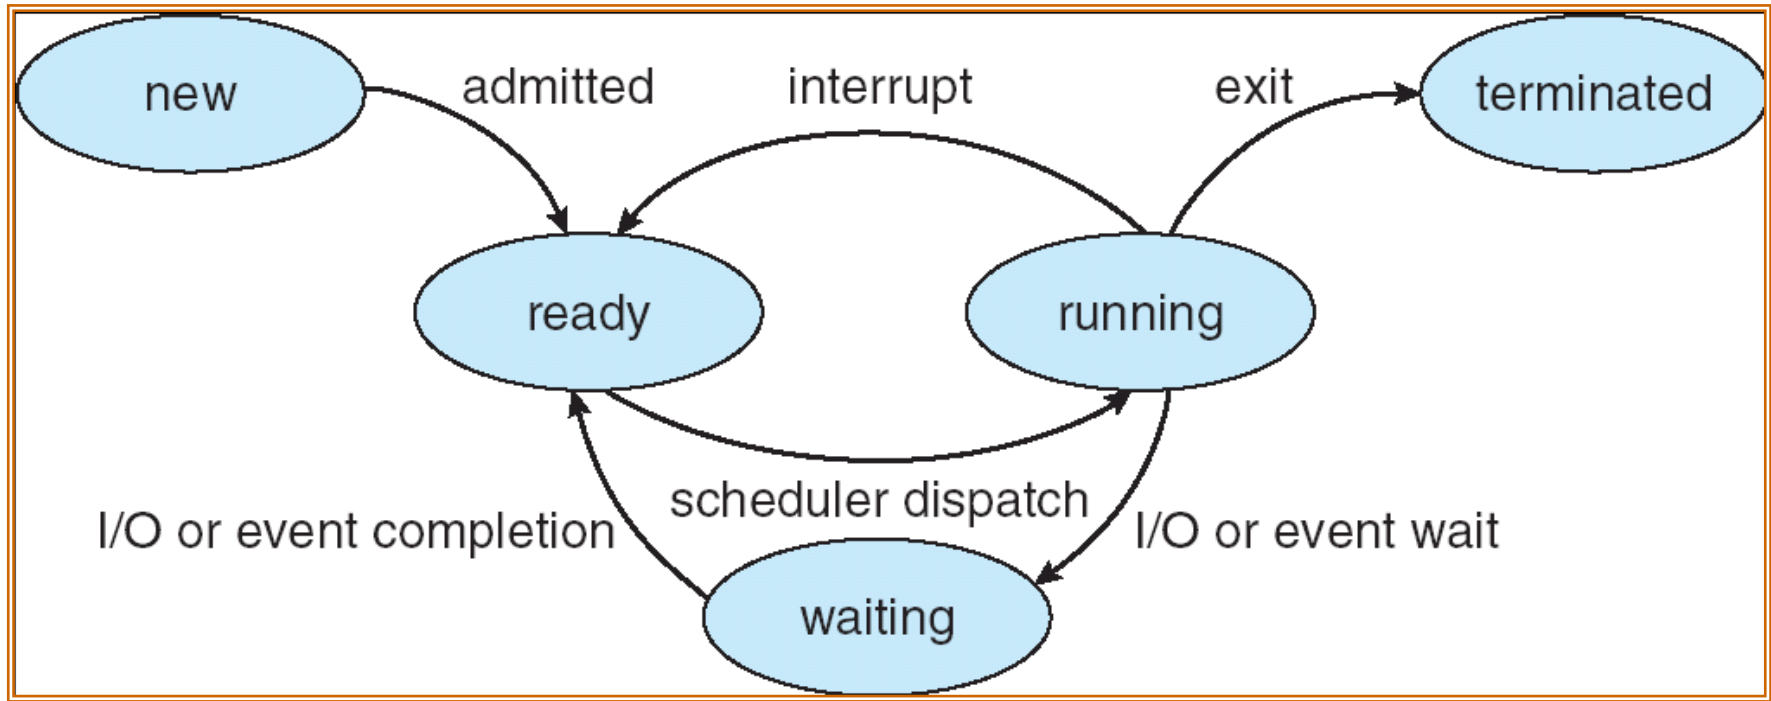


Process State

- As a process executes, it changes *state*
 - **new**: The process is **being** created
 - **running**: Instructions are being executed
 - **waiting**: The process is waiting for some event to occur
 - **ready**: The process is waiting to be assigned to a process
 - **terminated**: The process has finished execution



Diagram of Process State



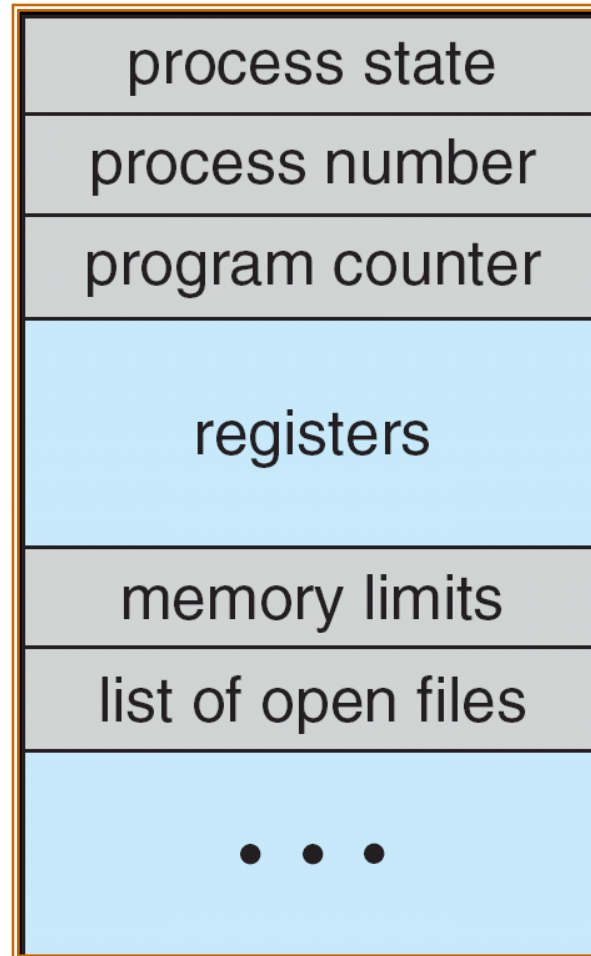
Process Control Block (PCB)

Information associated with each process

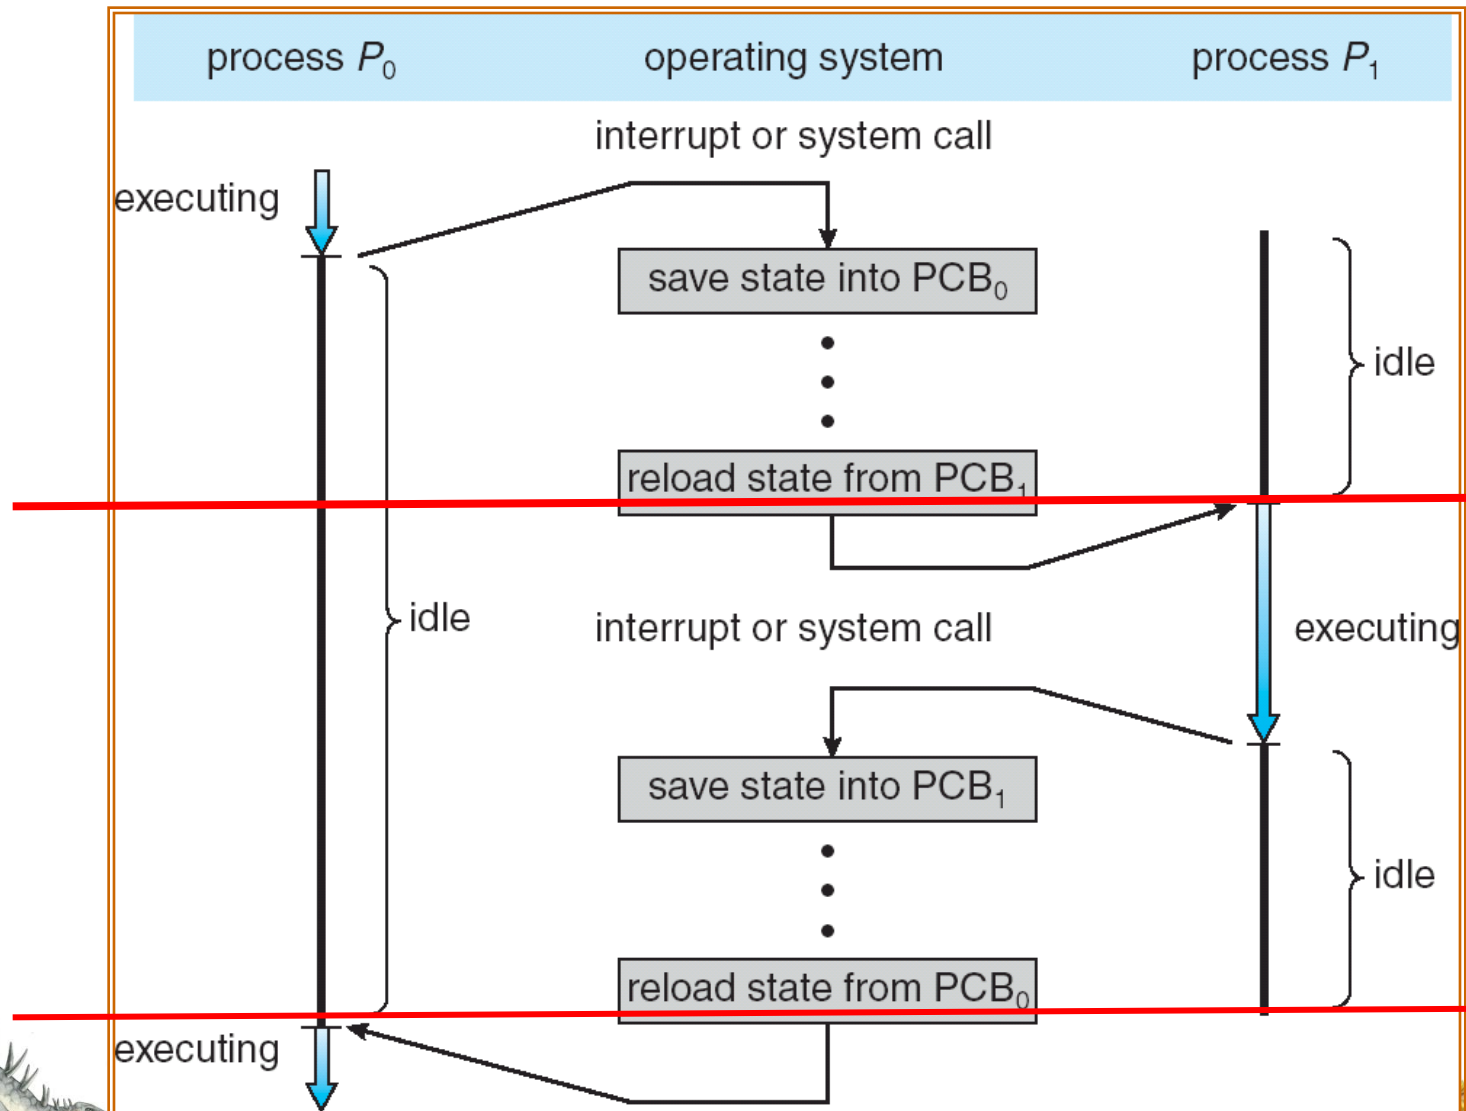
- ❑ Process state
- ❑ Program counter
- ❑ CPU registers
- ❑ CPU scheduling information
- ❑ Memory-management information
- ❑ Accounting information
- ❑ I/O status information



Process Control Block (PCB)



CPU Switch From Process to Process

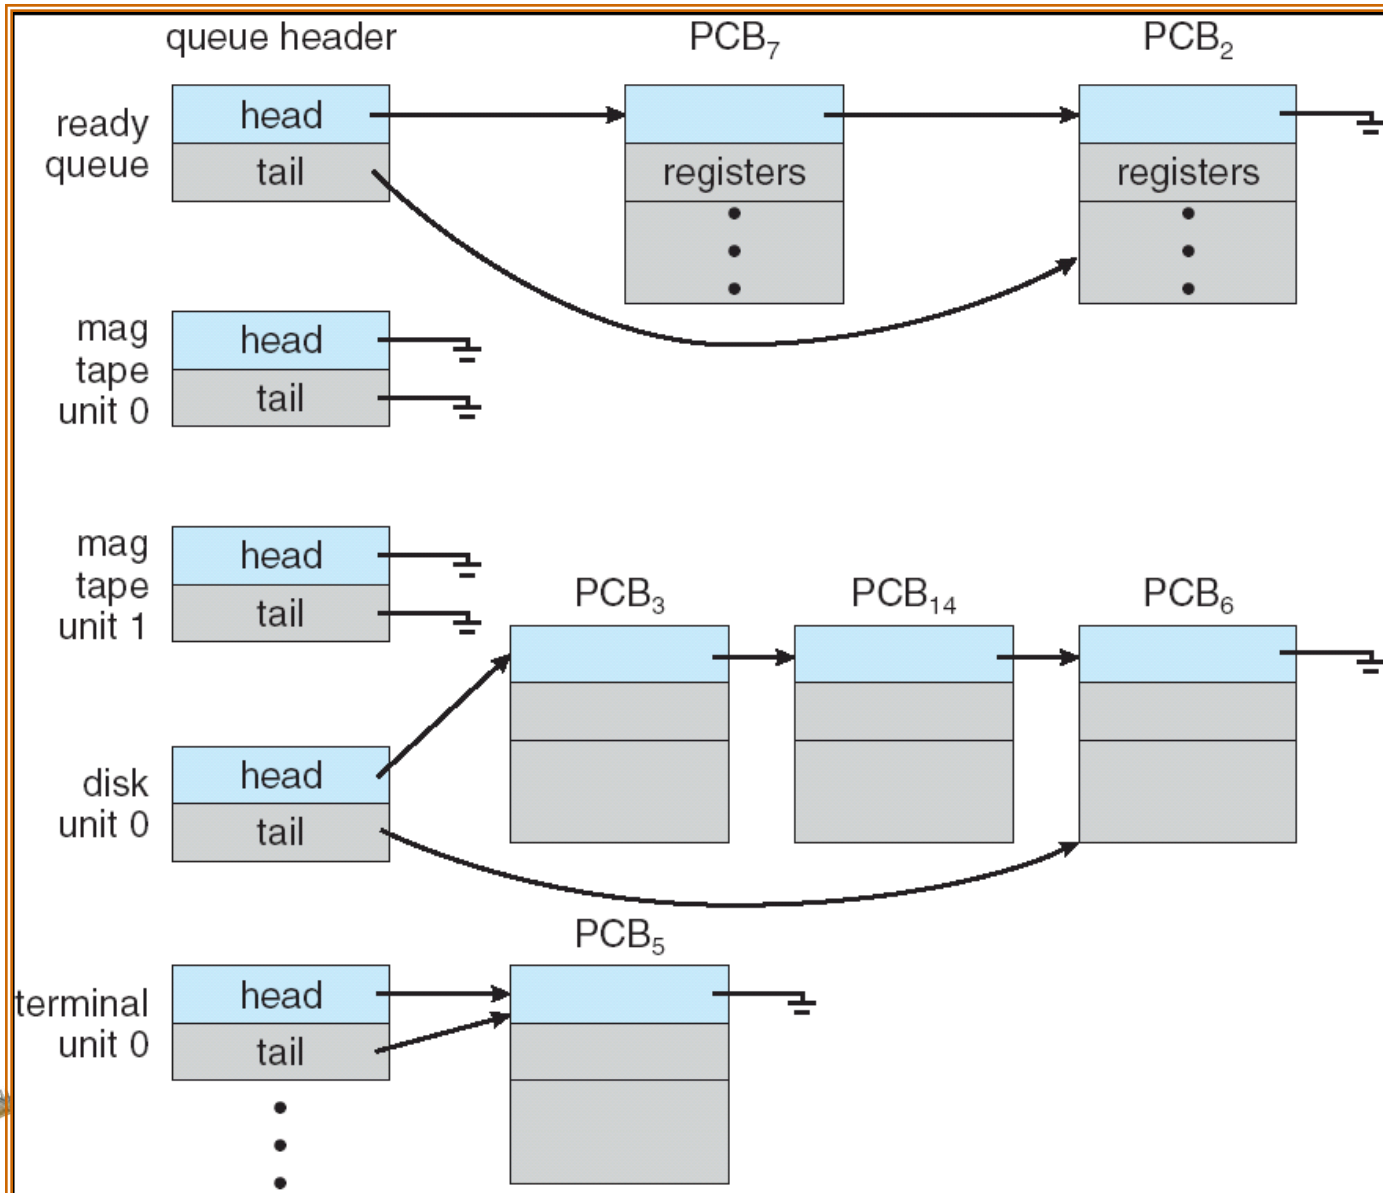


Process Scheduling Queues


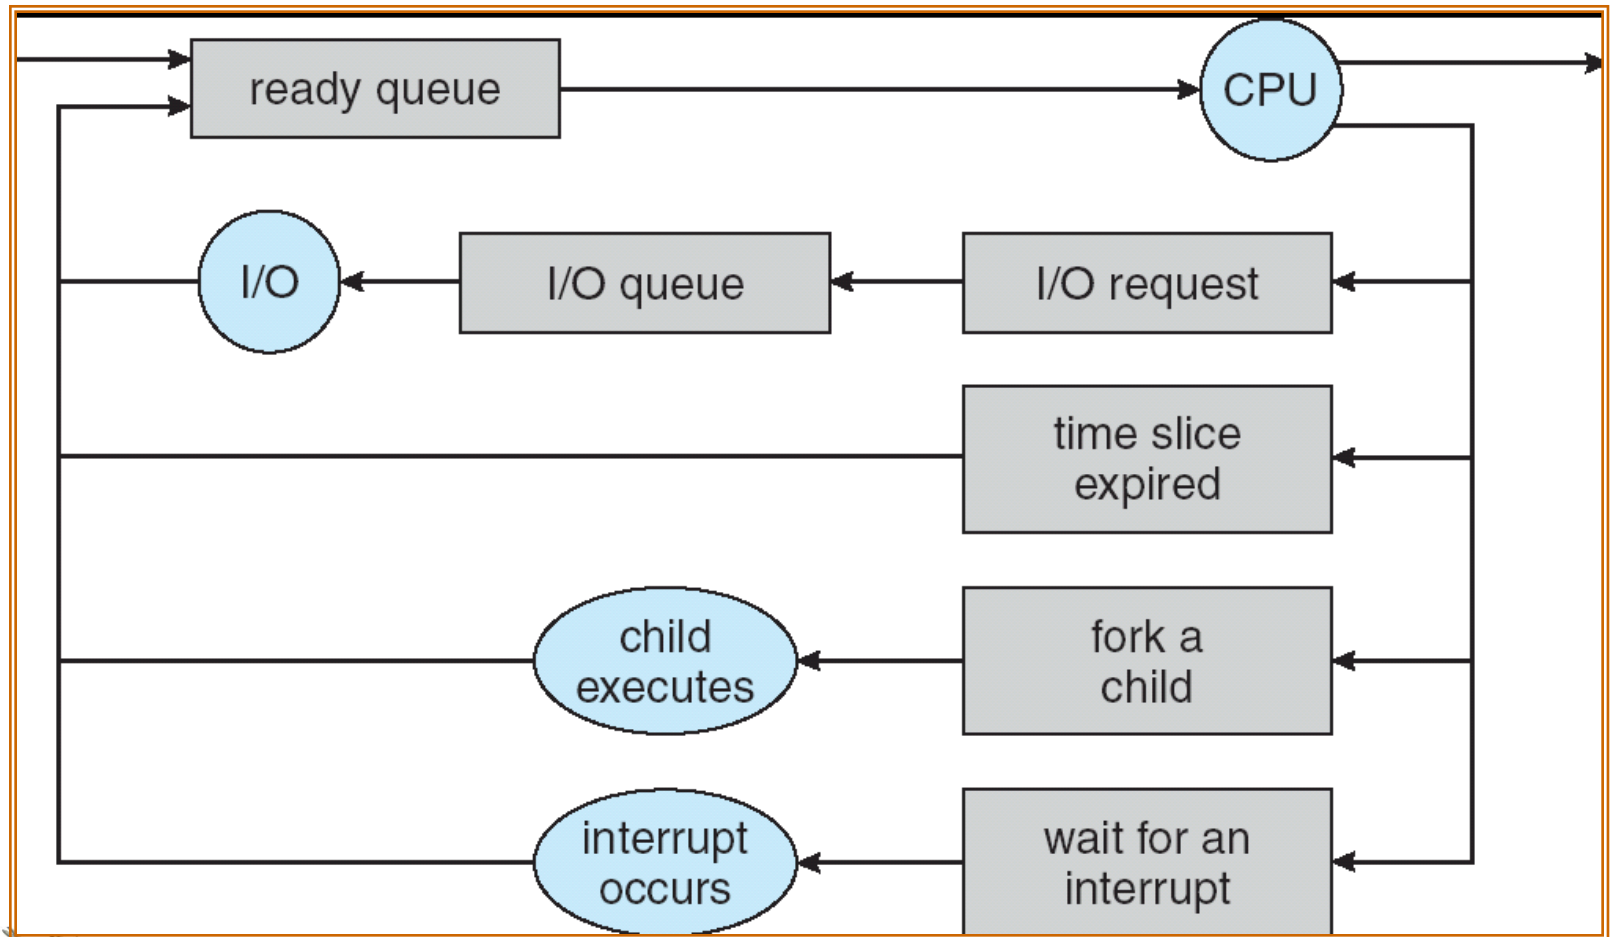
- ❑ **Job queue** – set of **all processes** in the system
- ❑ **Ready queue** – set of all processes **residing** in main memory, **ready** and **waiting** to execute
- ❑ **Device queues** – set of processes **waiting** for an I/O device
- ❑ Processes **migrate** among the various queues



Ready Queue And Various I/O Device Queues



Representation of Process Scheduling

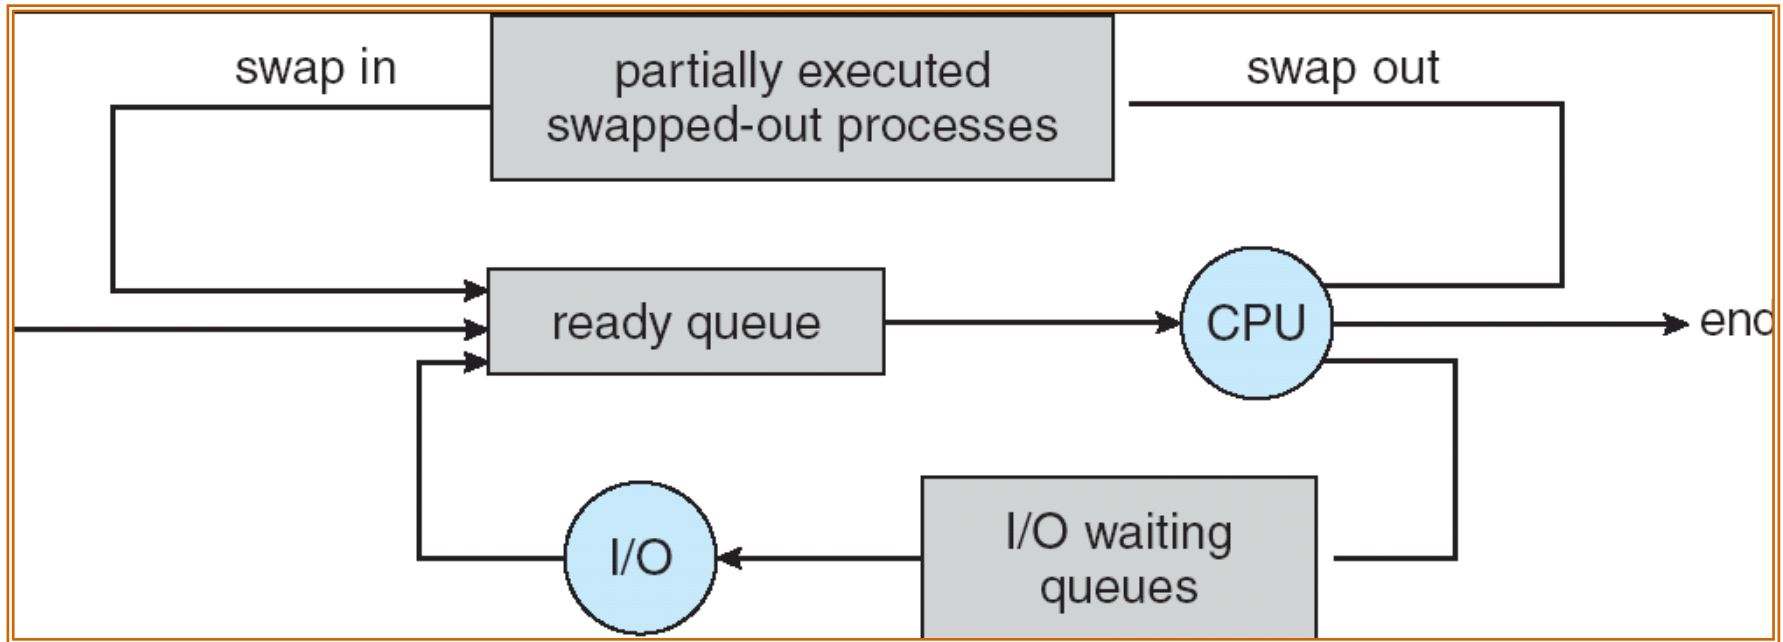


Schedulers

- ❑ **Long-term scheduler** (or job scheduler) – selects which processes should be **brought into the ready queue**
- ❑ **Short-term scheduler** (or CPU scheduler) – selects which process should be **executed next** and **allocates CPU**



Addition of **Medium Term** Scheduling



Schedulers (Cont.)

- ❑ Short-term scheduler is invoked **very frequently** (milliseconds) \Rightarrow (must be fast)
- ❑ Long-term scheduler is invoked **very infrequently** (seconds, minutes) \Rightarrow (may be slow)
- ❑ The long-term scheduler controls ***the degree of multiprogramming***
- ❑ Processes can be described as either:
 - **I/O-bound process** – spends more time doing I/O than computations, **many short CPU bursts**
 - **CPU-bound process** – spends more time doing computations; **few very long CPU bursts**



Context Switch

- ❑ When CPU **switches** to another process, the system must **save** the **state** of the old process and **load** the saved state for the new process
- ❑ Context-switch time is **overhead**; the system **does no useful work** while switching
- ❑ Time dependent on **hardware support**



Process Creation

- ❑ Parent process create children processes, which, in turn create other processes, **forming** a tree of processes
- ❑ Resource sharing
 - Parent and children share **all** resources
 - Children share **subset** of parent's resources
 - Parent and child share **no** resources
- ❑ Execution
 - Parent and children execute concurrently
 - Parent waits until children terminate

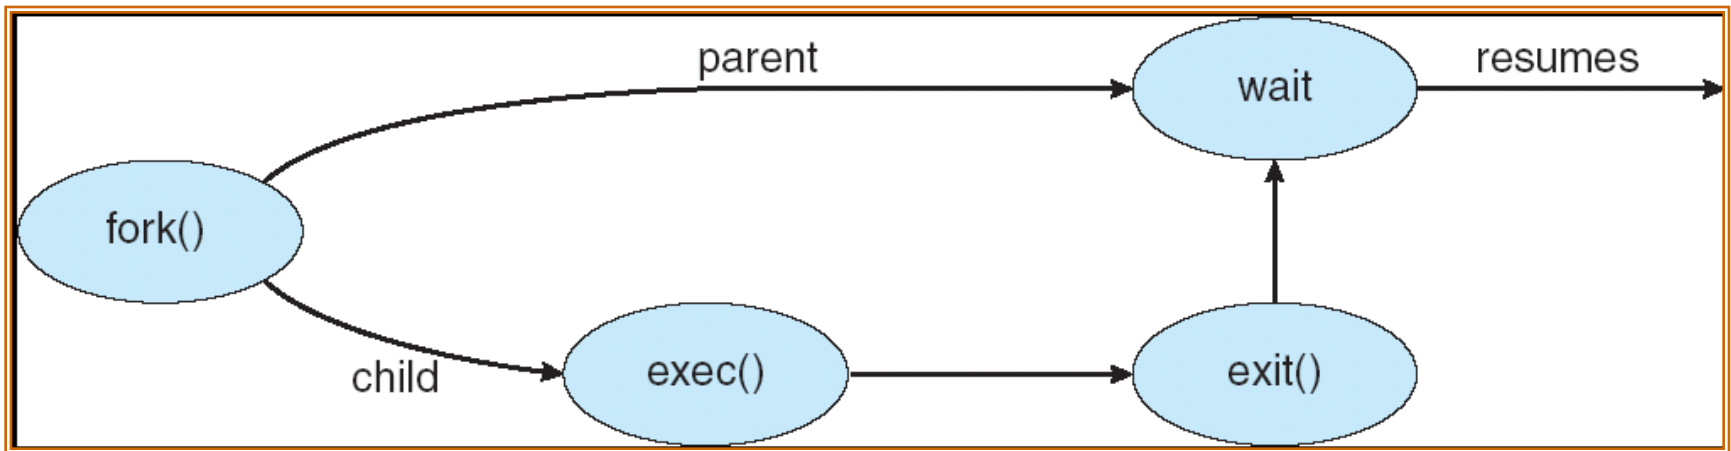


Process Creation (Cont.)

- ❑ Address space
 - Child **duplicate** of parent
 - Child has a program loaded into it
- ❑ UNIX examples
 - **fork** system call creates new process
 - **exec** system call used after a **fork** to replace the process' memory space with a new program



Process Creation



Process Creation in Win32

```
#include <stdio.h>
#include <windows.h>

int main(VOID)
{
    STARTUPINFO si;
    PROCESS_INFORMATION pi;

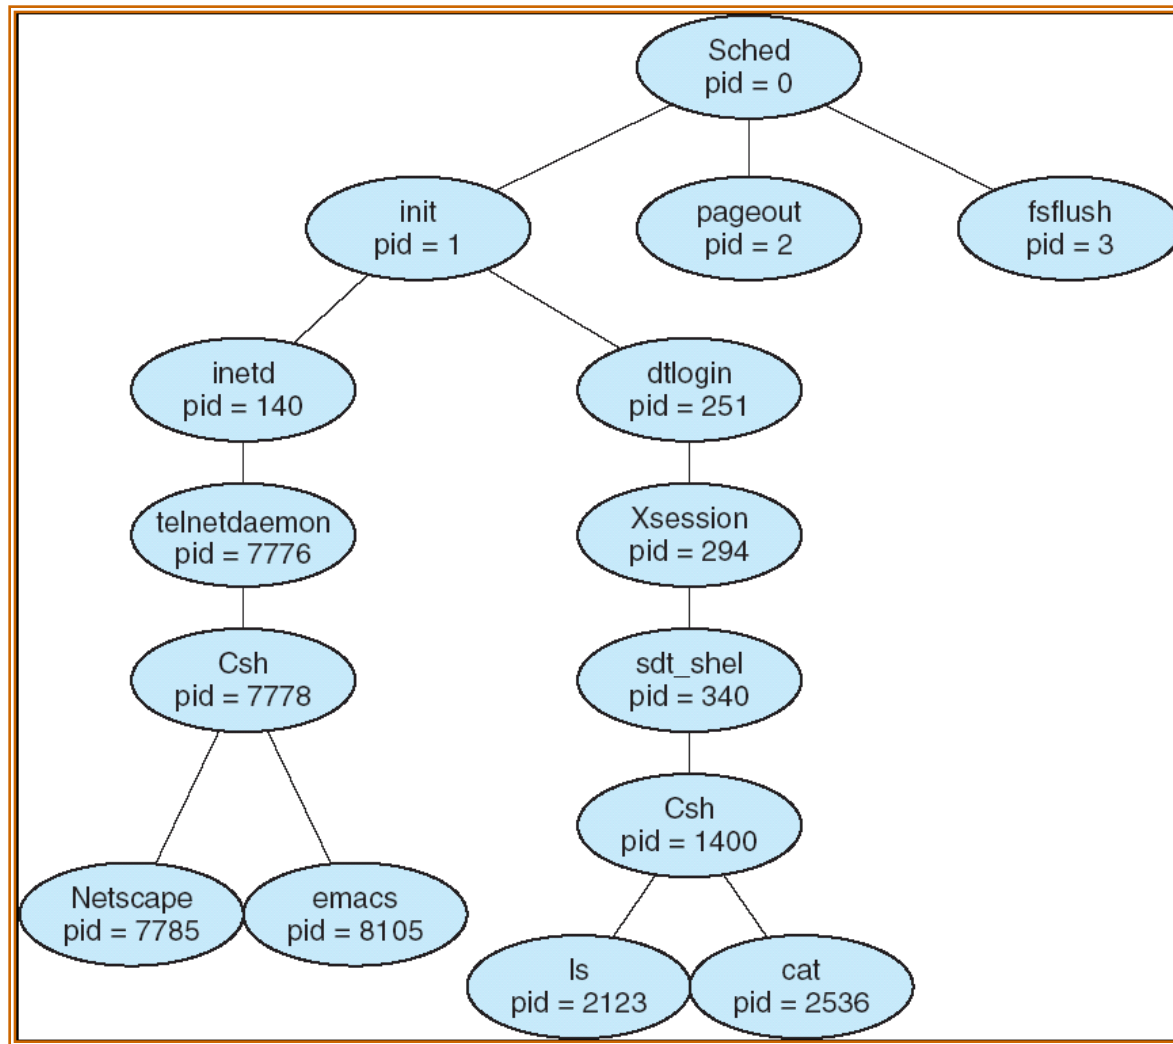
    // allocate memory
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));

    // create child process
    if (!CreateProcess(NULL, // use command line
        "C:\\WINDOWS\\system32\\mspaint.exe", // command line
        NULL, // don't inherit process handle
        NULL, // don't inherit thread handle
        FALSE, // disable handle inheritance
        0, // no creation flags
        NULL, // use parent's environment block
        NULL, // use parent's existing directory
        &si,
        &pi))
    {
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
    // parent will wait for the child to complete
    WaitForSingleObject(pi.hProcess, INFINITE);
    printf("Child Complete");

    // close handles
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
}
```



A tree of processes on a typical Solaris



Process Termination

- ❑ **Process executes last statement** and asks the operating system to delete it (**exit**)
 - Output data from child to parent (via **wait**)
 - Process' resources are **deallocated** by operating system
- ❑ Parent may terminate execution of **children processes** (**abort**)
 - Child has **exceeded** allocated resources
 - Task assigned to child is no longer required
 - If parent is exiting
 - ❑ Some operating system do not allow child to continue if its parent terminates
 - All children terminated - ***cascading termination***



Producer-Consumer Problem

- Paradigm for cooperating processes (processes that cooperate to achieve a specific task, for example, a task is initially broken into subtasks , each subtask is assigned to a process that should pass the result or input to another process and so on , until the task is done), *producer process* produces information that is consumed by a *consumer process*
 - *unbounded-buffer* places no practical limit on the size of the buffer (无限缓冲)
 - *bounded-buffer* assumes that there is a fixed buffer size (有限缓冲)



Bounded-Buffer – Shared-Memory Solution

□ Shared data

```
#define BUFFER_SIZE 10
```

```
typedef struct {
```


```
    . . .
```

```
} item;
```

```
item buffer[BUFFER_SIZE]; // 有限缓冲
```

```
int in = 0; // 指向缓冲中下一个空位
```

```
int out = 0; // 指向缓冲中第一个满位
```



Bounded-Buffer – Insert() Method

```
while (true) {  
    /* Produce an item */  
    while (( (in + 1) % BUFFER SIZE) == out)  
        ; /* do nothing -- no free buffers */  
    buffer[in] = item;  
    in = (in + 1) % BUFFER SIZE;  
}
```



Bounded Buffer – Remove() Method

```
while (true) {  
    while (in == out)  
        ; // do nothing -- nothing to consume  
  
    // remove an item from the buffer  
    item = buffer[out];  
    out = (out + 1) % BUFFER SIZE;  
    return item;  
}
```




Bounded-Buffer – Shared-Memory Solution

- Solution is correct, but can only use **BUFFER_SIZE-1** elements.





Interprocess Communication

Message Passing



Shared Memory



Message Passing & Shared Memory

- ❑ Message system – processes communicate with each other without **resorting** to shared variables
- ❑ Message passing facility provides two operations:
 - **send**(*message*) – message size **fixed** or **variable**
 - **receive**(*message*)
- ❑ If P and Q wish to communicate, they need to:
 - establish a **communication link** between them
 - exchange messages via **send/receive**
- ❑ Implementation of communication link
 - physical (e.g., shared memory, hardware bus)
 - logical (e.g., logical properties)
 - ❑ 直接或间接通信、同步或异步通信



Implementation Questions

- ❑ How are links established?
- ❑ Can a link be associated with more than two processes?
- ❑ How many links can there be between every pair of communicating processes?
- ❑ What is the capacity of a link?
- ❑ Is the size of a message that the link can accommodate fixed or variable?
- ❑ Is a link unidirectional or bi-directional?



Direct Communication

- ❑ Processes must name each other explicitly:
 - **send** (P , *message*) – send a message **to process P**
 - **receive**(Q , *message*) – receive a message **from process Q**
- ❑ Properties of communication link
 - Links are established **automatically**
 - A link is associated with exactly **one pair of communicating processes**
 - Between each pair there exists exactly one link
 - The link may be **unidirectional**, but is usually bi-directional



Indirect Communication

- ❑ Messages are directed and received from **mailboxes** (also referred to as ports)
 - Each mailbox has a **unique id**
 - Processes can communicate only if they share a mailbox
- ❑ Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with **many processes**
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional



Indirect Communication

- Operations
 - create a new mailbox
 - send and receive messages through mailbox
 - destroy a mailbox
- Primitives are defined as:
 - send**(*A, message*) – send a message to mailbox *A*
 - receive**(*A, message*) – receive a message from mailbox *A*



Indirect Communication

- ❑ Mailbox sharing
 - P_1 , P_2 , and P_3 share mailbox A
 - P_1 sends; P_2 and P_3 receive
 - Who gets the message?
- ❑ Solutions
 - Allow a link to be associated with at most two processes
 - Allow **only one process** at a time to execute a **receive** operation
 - Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.



Synchronization & Asynchronous

- ❑ Message passing may be either blocking or non-blocking
- ❑ **Blocking** is considered **synchronous**
 - **Blocking send** has the sender **block** until the message is received
 - **Blocking receive** has the receiver **block** until a message is available
- ❑ **Non-blocking** is considered **asynchronous**
 - **Non-blocking send** has the sender send the message and **continue**
 - **Non-blocking receive** has the receiver receive a valid message or null



Buffering

- Queue of messages attached to the link;
implemented in one of three ways
 1. **Zero capacity** – 0 messages
Sender must wait for receiver (**rendezvous**)
 2. **Bounded capacity** – **finite** length of n messages
Sender must wait if link full
 3. **Unbounded capacity** – **infinite** length
Sender never waits



Client-Server Communication

- Sockets
- Remote Procedure Calls
- Remote Method Invocation (Java)

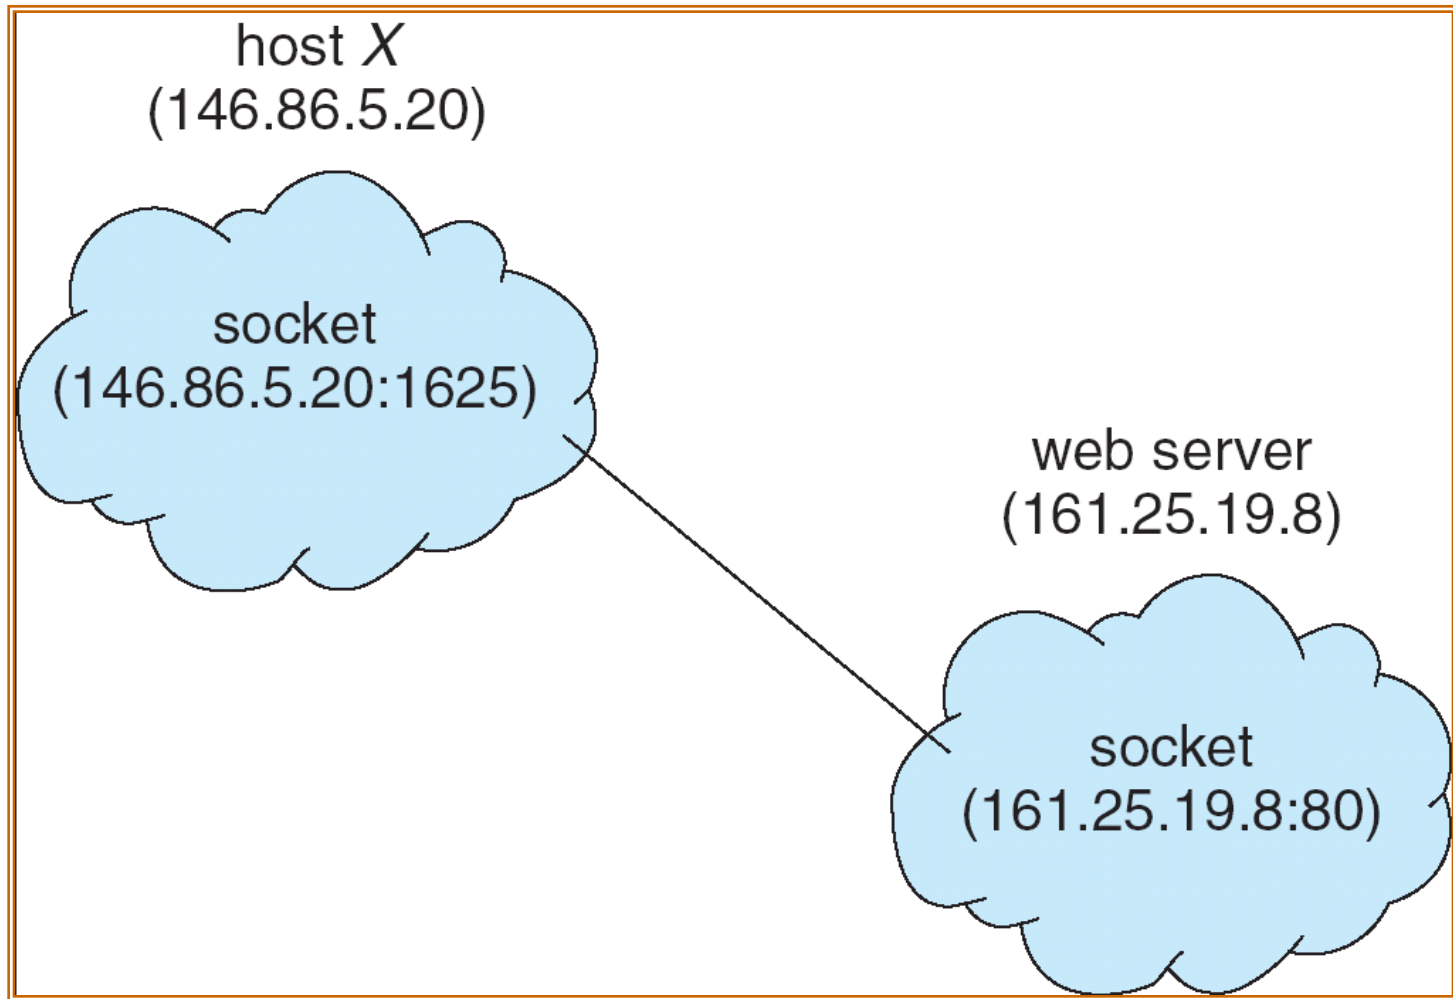


Sockets

- ❑ A socket is defined as an *endpoint for communication*
- ❑ Concatenation of IP address and port
- ❑ The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
- ❑ Communication consists between a pair of sockets



Socket Communication



Remote Procedure Calls

- ❑ Remote procedure call (RPC) abstracts procedure calls between processes on networked systems.
- ❑ **Stubs** – client-side proxy for the actual procedure on the server.
- ❑ The client-side stub locates the server and *marshalls* the parameters.
- ❑ The server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server.



End of Chapter 3

