

Chapter 4: Threads



Chapter 4: Threads

- Overview
- **Multithreading** Models
- Threading Issues
- Java Threads
- Pthreads



Process/Thread

- ❑ **Resource ownership** - process includes a virtual address space to hold the process image
- ❑ **Scheduling/execution** - follows an execution path that may be interleaved with other processes
- ❑ These two characteristics are treated independently by the operating system
 - **Process**: Ownership of memory, files, other resources
 - ❑ 资源分配的基本单元
 - **Thread**: Unit of execution we use to dispatch
 - ❑ 指令执行的基本单元

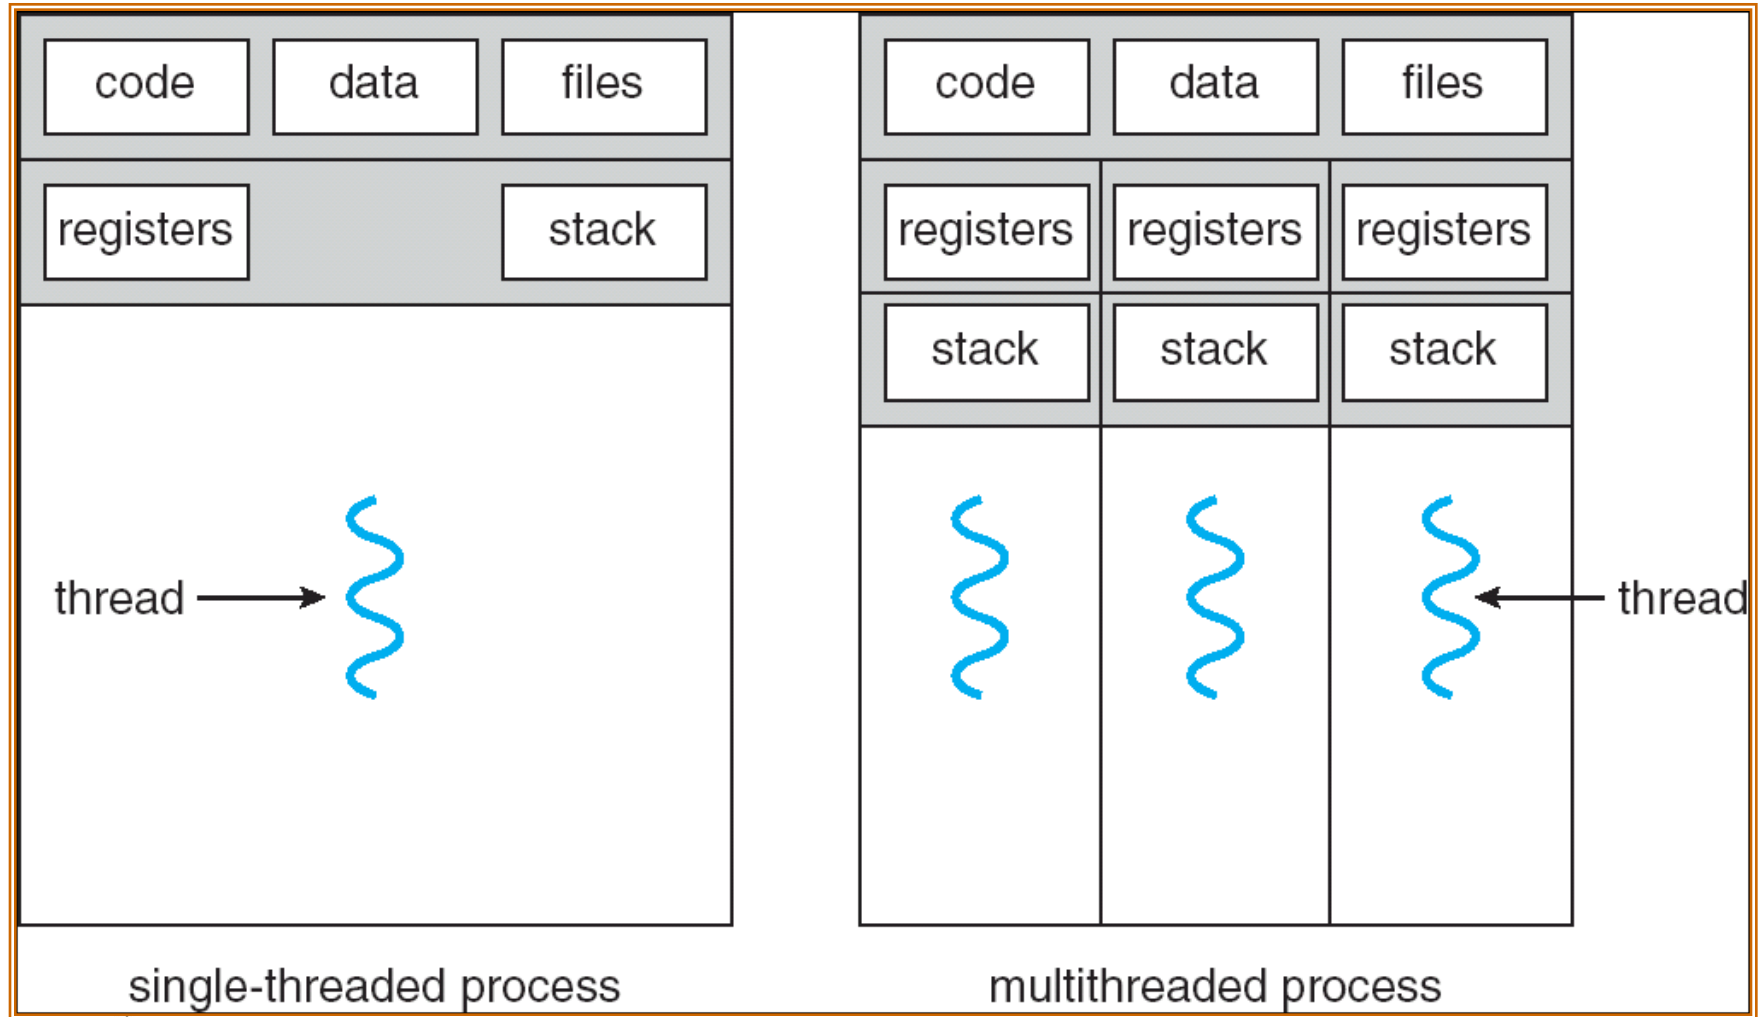


Thread Control Block

- ❑ Each Thread has a **Thread Control Block** (TCB)
 - Execution State: CPU registers, program counter, pointer to **stack**
 - Scheduling Info: State, priority, CPU time
 - Accounting Info
 - Various Pointers (for implementing scheduling queues)
 - Pointer to enclosing process (PCB)




Single and Multithreaded Processes




Single and Multithreaded Processes

Single-Threaded Process Model



Multithreaded Process Model



Benefits

- ❑ Responsiveness (响应速度快)
- ❑ Resource Sharing (共享进程的资源)
- ❑ Economy (经济)
- ❑ Utilization of MP (多处理器) Architectures



User and Kernel Threads

- ❑ **User threads** - Thread management done by user-level threads library.
 - Three primary thread libraries:
 - ❑ POSIX Pthreads
 - ❑ Win32 threads
 - ❑ Java threads
- ❑ **Kernel threads** - Threads directly supported by the kernel.
 - Downside of kernel threads: a bit expensive



Kernel Threads

Examples (多进程, 多线程):

- Windows XP/2000
- Solaris
- Linux
- Tru64 UNIX
- Mac OS X



Multithreading Models

Mapping user threads to **kernel threads**:

- ❑ **Many-to-One**
- ❑ **One-to-One**
- ❑ **Many-to-Many**




Many-to-One

- ❑ Many user-level threads **mapped** to single kernel thread
- ❑ Examples:
 - Solaris Green Threads
 - GNU Portable Threads



Many-to-One Model




One-to-One

- **Each** user-level thread maps to kernel thread
- Examples
 - Windows NT/XP/2000
 - Linux
 - Solaris 9 and later



One-to-one Model




Many-to-Many Model

- ❑ Allows **many** user level threads to be mapped to **many** kernel threads
- ❑ Allows the operating system to create a **sufficient number** of kernel threads
- ❑ Solaris prior to version 9
- ❑ Windows NT/2000 with the *ThreadFiber* package



Many-to-Many Model




Two-level Model

- Similar to M:M, except that it allows a user thread to be **bound** to kernel thread
- Examples
 - IRIX
 - HP-UX
 - Tru64 UNIX
 - Solaris 8 and earlier



Two-level Model



Java Threads

- ❑ Java threads are managed by the **JVM**
- ❑ Java threads may be created by:
 - Implementing the Runnable interface

```
public interface Runnable
{
    public abstract void run();
}
```



Java Threads - Example Program

```
class MutableInteger
{
    private int value;
    public int getValue() {
        return value;
    }
    public void setValue(int value) {
        this.value = value;
    }
}

class Summation implements Runnable
{
    private int upper;
    private MutableInteger sumValue;
    public Summation(int upper, MutableInteger sumValue) {
        this.upper = upper;
        this.sumValue = sumValue;
    }
    public void run() {
        int sum = 0;
        for (int i = 0; i <= upper; i++)
            sum += i;
        sumValue.setValue(sum);
    }
}
```



Java Threads - Example Program

```
public class Driver
{
    public static void main(String[] args) {
        if (args.length > 0) {
            if (Integer.parseInt(args[0]) < 0)
                System.err.println(args[0] + " must be >= 0.");
            else {
                // create the object to be shared
                MutableInteger sum = new MutableInteger();
                int upper = Integer.parseInt(args[0]);
                Thread thrd = new Thread(new Summation(upper, sum));
                thrd.start();
                try {
                    thrd.join();
                    System.out.println
                        ("The sum of "+upper+" is "+sum.getValue());
                } catch (InterruptedException ie) { }
            }
        }
        else
            System.err.println("Usage: Summation <integer value>");
    }
}
```



Thread Pools

- ❑ Create **a number of threads** in a pool where they **await** work
- ❑ Advantages:
 - Usually **slightly faster to service** a request with an existing thread than create a new thread
 - Allows the number of threads in the application(s) to be bound to the size of the pool



Thread Pools

□ Java provides 3 thread pool architectures:

1. **Single thread executor** - pool of size 1.

- `static ExecutorService newSingleThreadExecutor()`

2. **Fixed thread executor** - pool of fixed size.

- `static ExecutorService newFixedThreadPool(int nThreads)`


3. **Cached thread pool** - pool of **unbounded** size

- `static ExecutorService newCachedThreadPool()`



Scheduler Activations

- ❑ Both M:M and Two-level models require communication to maintain the **appropriate number of kernel threads** allocated to the application
- ❑ Scheduler activations provide **upcalls** - a communication mechanism from the kernel to the thread library
- ❑ This communication allows an application to maintain the **correct number kernel** threads



Pthreads

- ❑ A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- ❑ API specifies behavior of the thread library, implementation is up to development of the library
- ❑ Common in UNIX operating systems (Solaris, Linux, Mac OS X)



End of Chapter 4

