

Chapter 8: Main Memory



课前习题

桌上有一个空盘，允许存放一只水果。爸爸可以向盘中放苹果，也可向盘中放桔子，儿子专等吃盘中的桔子，女儿专等吃盘中的苹果。规定当盘空时一次只能放一只水果供吃者使用，请用P、V原语实现爸爸、儿子、女儿三个并发进程的同步。**提示：生产者、消费者问题的变形**



Binary_semaphore ore S =1;	Father() {	son() {	daughter() {
Binary_semaphore ore Sa =0;	While(1) {	While(1) {	While(1) {
Binary_semaphore ore So =0;	P(S);	P(So);	P(Sa);
Main() {	放水果到盘中	取桔子	取苹果
begin	;	V(S)	V(S)
Father()	If(放的是桔	吃桔子	吃苹果
Son()	子)	}	}
Daughter();	V(So))	}	}
end;	Else	}	}
}	}	}	}
	}		



课前习题

假定某计算机系统有R1和R2两类可再使用资源，其中R1有两个单位，R2有一个单位，它们被进程P1和P2所共享，且已知两个进程均以下列顺序使用两类资源：

—申请R1—申请R2—申请R1—释放R1—释放R2—释放R1—

试求出系统运行过程中可能到达的死锁点，并画出死锁点的资源分配图。



课前习题

假定某计算机系统有R1和R2两类资源，其中R1有两个单位，R2有一个单位，它们被进程P1和P2所共享，且已知两个进程均以下列顺序使用两类资源：

—申请R1—申请R2—申请R1—释放R1—释放R2—释放R1—

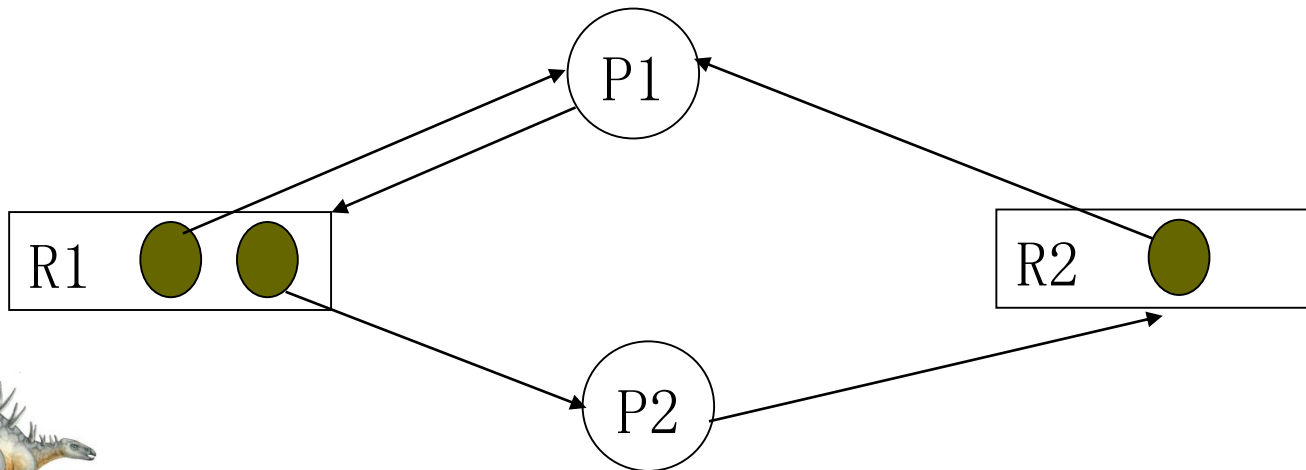
试求出系统运行过程中可能到达的死锁点，并画出死锁点的资源分配图。

本题中，当两个进程都执行完第一步后，即进程P1和P2都申请到了一个R1类资源时，系统进入不‘安全状态。随着两个进程的向前推进，无论哪个进程执行完第二步，系统都将进入死锁状态。

可能的到达的死锁点是：

P1占有一个单位R1及一个单位R2，P2占有一个单位R1。

P2占有一个单位R1及一个单位R2，P1占有一个单位R1。



Chapter 8: Memory Management

- ❑ Background
- ❑ Contiguous Memory Allocation
- ❑ Paging
- ❑ Structure of the Page Table
- ❑ Segmentation
- ❑ Example: The Intel Pentium



Objectives

- ❑ To provide a detailed description of various ways of **organizing** memory hardware
- ❑ To discuss various memory-management techniques, including paging and segmentation
- ❑ To provide a detailed description of the Intel Pentium, which supports both pure segmentation and segmentation with paging



Background

- ❑ Program **must** be brought (from disk) into **memory** and placed within a process for it to be run
- ❑ **Main memory** and **registers** are only storage CPU can access **directly**
- ❑ Register access **in one CPU clock** (or less)
- ❑ Main memory **can take many cycles**
- ❑ **Cache** sits **between** main memory and CPU registers
- ❑ Protection of memory required to ensure correct operation



Logical vs. Physical Address Space

- ❑ The concept of a logical address space that is bound to a separate **physical address space** is central to proper memory management
 - **Logical address** – generated by the CPU; also referred to as **virtual address**
 - **Physical address** – address seen by the memory unit
- ❑ Logical and physical addresses **are the same** in **compile-time and load-time address-binding schemes**; logical (virtual) and physical addresses **differ** in **execution-time address-binding scheme**



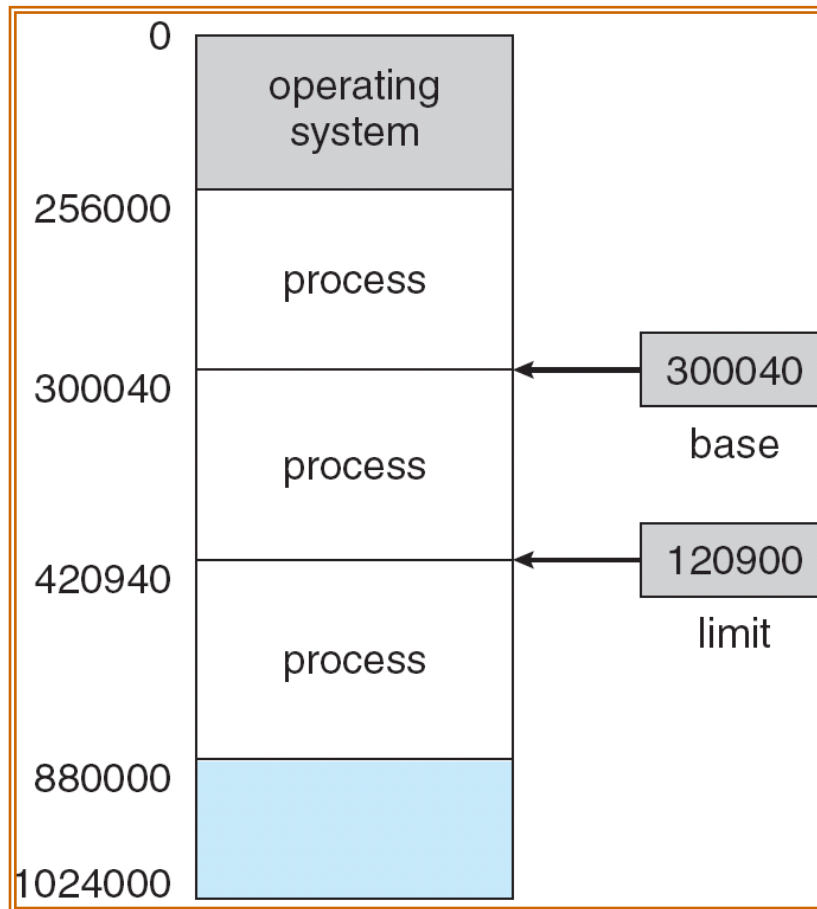
Memory-Management Unit (MMU)

- ❑ Hardware device that **maps virtual to physical address**
- ❑ In MMU scheme, the value in the relocation register is added to every address generated by a user process at the time it is sent to memory
- ❑ The user program deals with **logical addresses**; it never sees the *real*/physical addresses

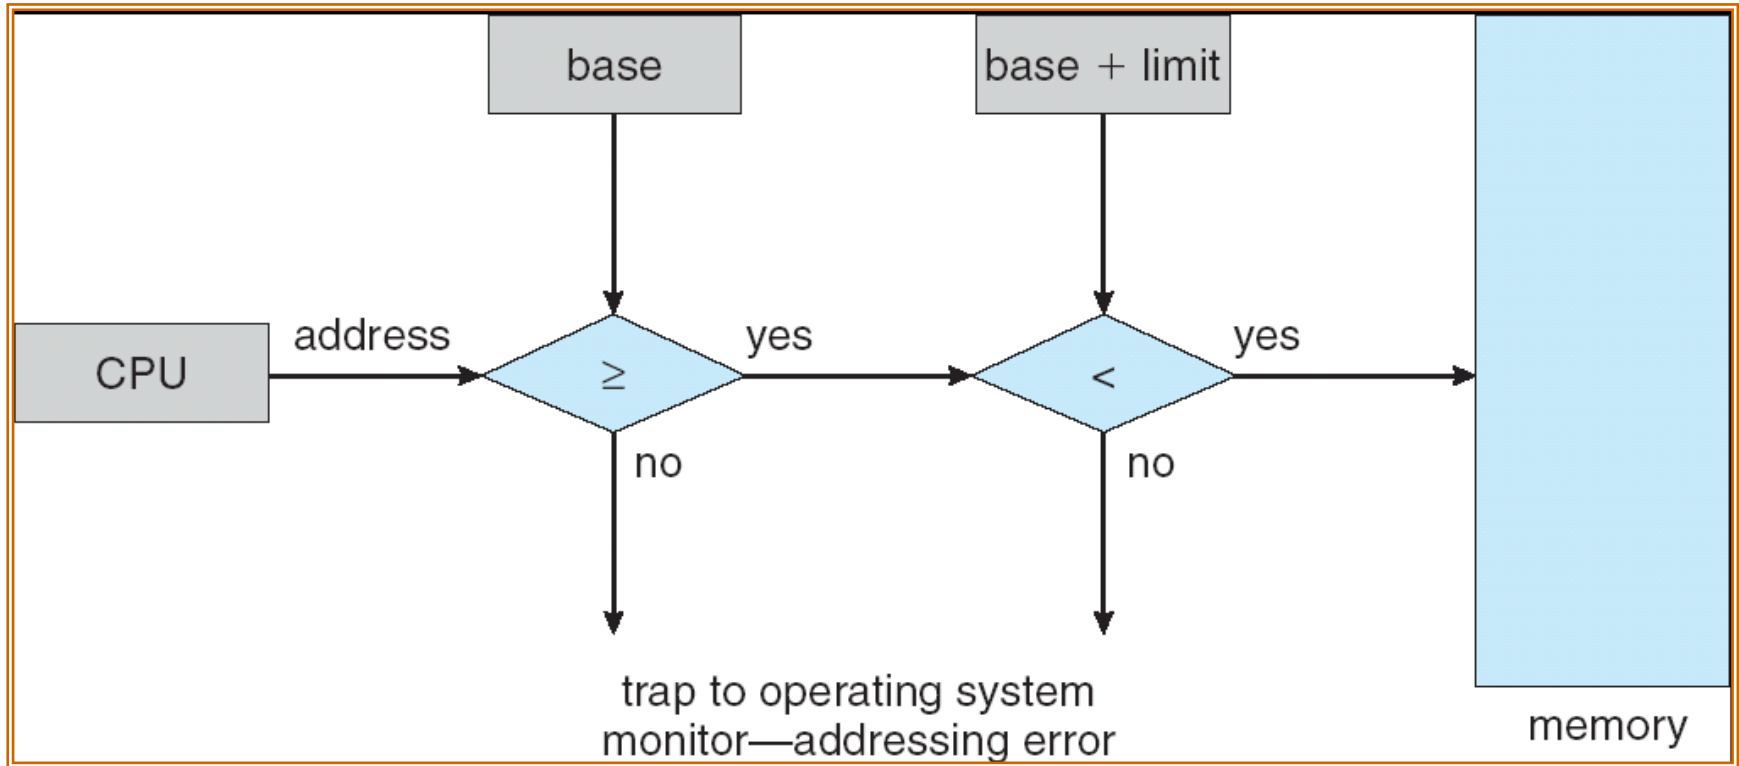


Base and Limit Registers

- A pair of **base** and **limit** registers define the logical address space



HW address protection with base and limit registers



Contiguous Allocation

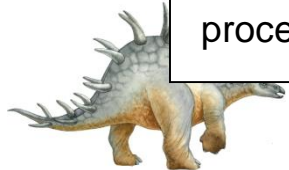
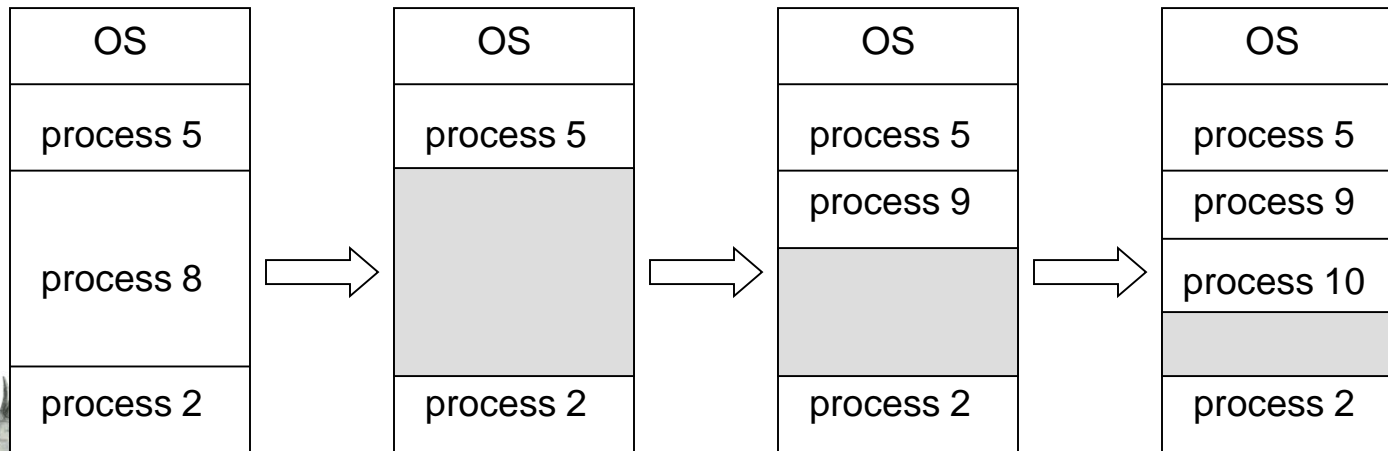
- ❑ Main memory usually into two partitions:
 - Resident operating system, **usually** held in **low memory** with interrupt vector
 - User processes then held in **high memory**

- ❑ Relocation registers used to **protect user processes from each other**, and from changing operating-system code and data
 - Base register contains value of **smallest physical address**
 - Limit register contains range of logical addresses – each logical address must be less than the limit register
 - MMU maps logical address *dynamically*



Contiguous Allocation (Cont.)

- ❑ Multiple-partition allocation
 - **Hole** – block of available memory; holes of various size are scattered throughout memory
 - When a process arrives, it is allocated memory from a **hole large enough to accommodate it**
 - Operating system maintains information about:
 - a) allocated partitions b) free partitions (hole)



Dynamic Storage-Allocation Problem

How to satisfy a request of size n from a list of free holes

- ❑ **First-fit:** Allocate the *first* hole that is big enough
- ❑ **Best-fit:** Allocate the *smallest* hole that is big enough; must search entire list, unless ordered by size
 - Produces the smallest leftover hole
- ❑ **Worst-fit:** Allocate the *largest* hole; must also search entire list
 - Produces the largest leftover hole

First-fit and best-fit better than worst-fit in terms of **speed and storage utilization**



Fragmentation

- ❑ **External Fragmentation** – total memory space exists to satisfy a request, but it **is not contiguous**
- ❑ **Internal Fragmentation** – allocated memory may be slightly **larger** than requested memory; this size difference is memory internal to a partition, but not being used
- ❑ Reduce external fragmentation by **compaction**
 - Shuffle memory contents to place all free memory **together** in one large block
 - Compaction is possible **only if relocation is dynamic**, and is done at execution time



分页技术

大小不等的固定分区和大小可变的分区效率低下，会产生内部或外部碎片。若将主存划分为**固定大小块**，并且**块较小**，每个进程也被分为同样大小的小块，可以将进程的块（称为页）指定到内存中的可用块（称为帧或页帧），最后一页**可能形成内部碎片**，没有外部碎片。



Paging

- ❑ Logical address space of a process can be **noncontiguous**; process is allocated physical memory whenever the latter is available
- ❑ Divide physical memory **into fixed-sized blocks** called **frames** (size is power of 2, between 512 bytes and 8,192 bytes)
- ❑ Divide **logical memory** into blocks of same size called **pages**
- ❑ Keep **track** of **all free frames**
- ❑ To run a program of size **n** pages, need to find **n** free frames and load program
- ❑ Set up a page table **to translate logical to physical addresses**
- ❑ Internal fragmentation



分页技术

Frame number	Main memory
0	
1	
2	
3	
4	
5	
6	
7	
8	
9	
10	
11	
12	
13	
14	

(a) Fifteen Available Frames

	Main memory
0	A.0
1	A.1
2	A.2
3	A.3
4	
5	
6	
7	
8	
9	
10	
11	
12	
13	
14	

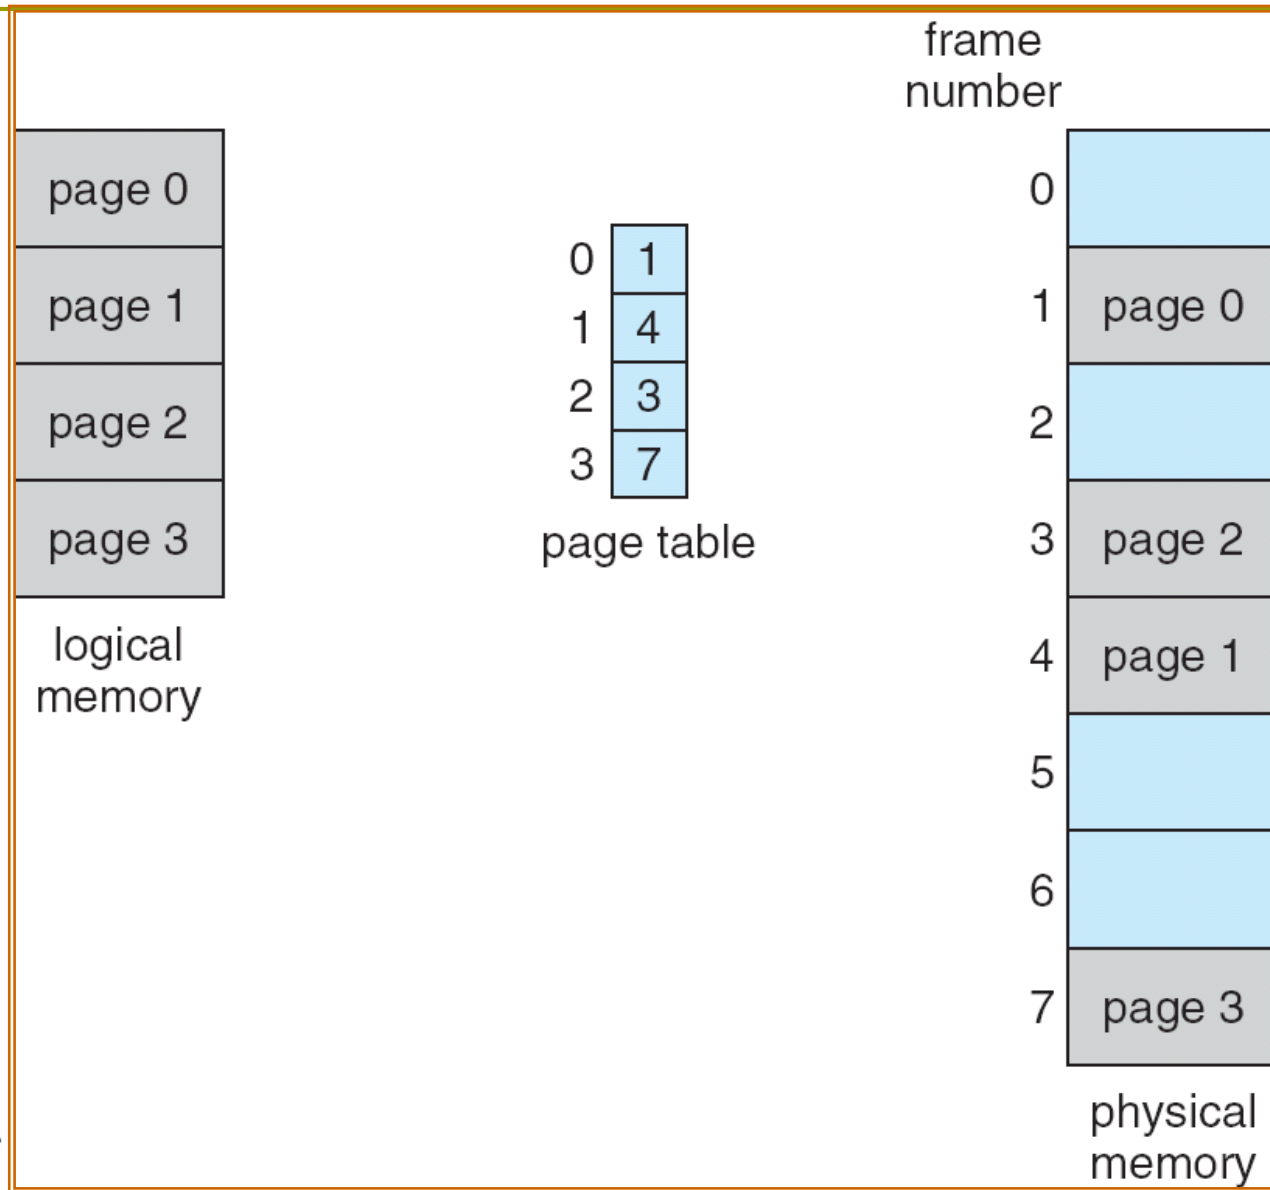
(b) Load Process A

	Main memory
0	A.0
1	A.1
2	A.2
3	A.3
4	B.0
5	B.1
6	B.2
7	
8	
9	
10	
11	
12	
13	
14	

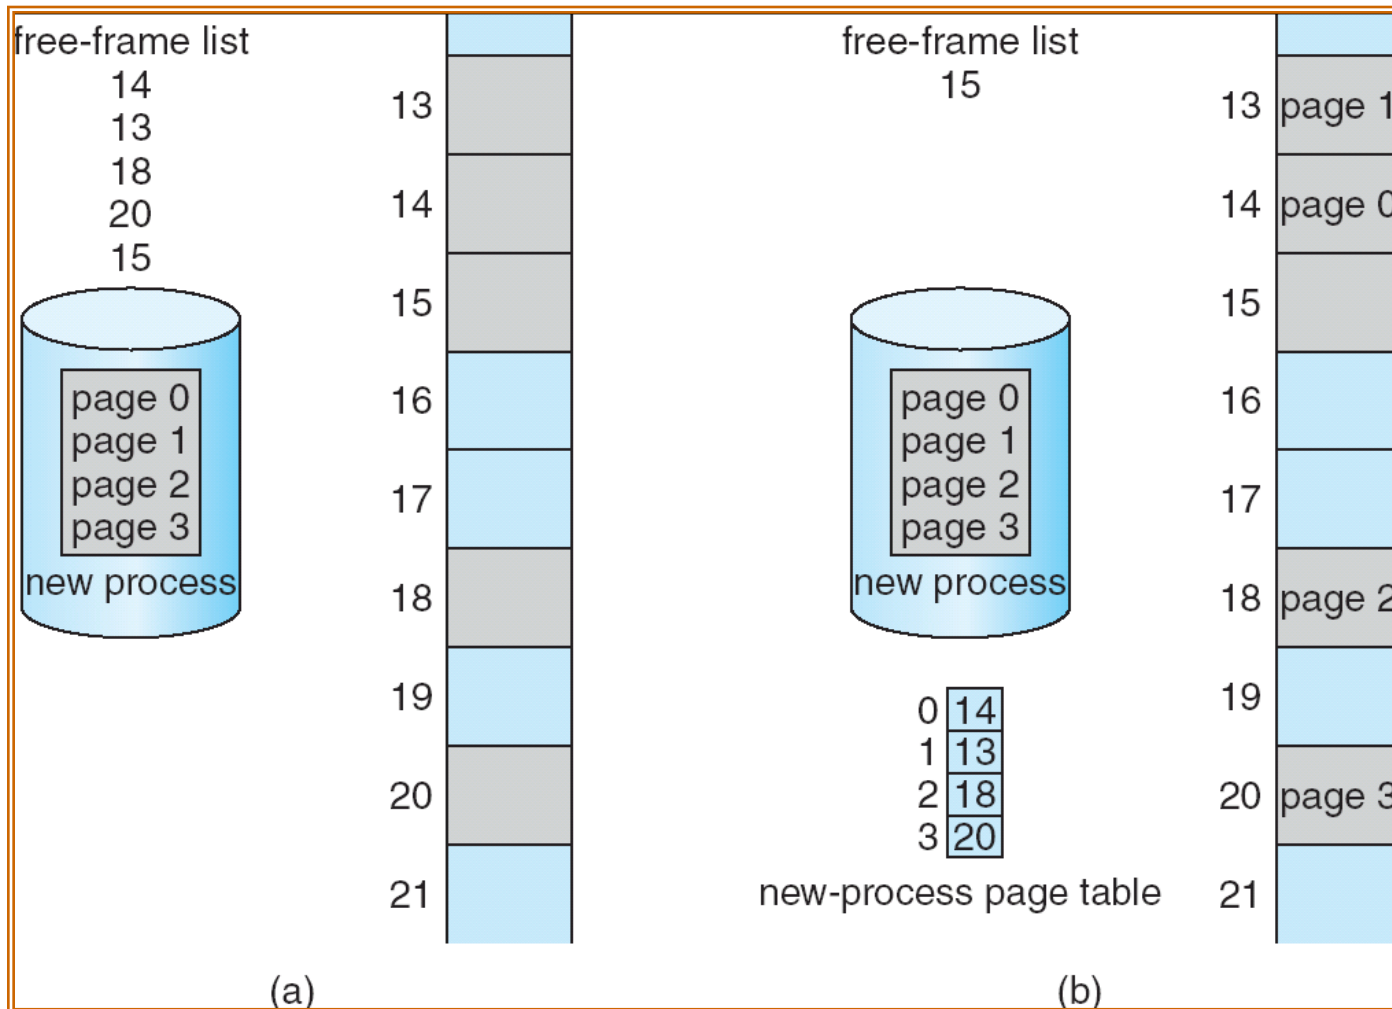
(c) Load Process B



Paging Model of Logical and Physical Memory



Free Frames



Before allocation

After allocation



分页技术

Main memory

0	A.0
1	A.1
2	A.2
3	A.3
4	B.0
5	B.1
6	B.2
7	C.0
8	C.1
9	C.2
10	C.3
11	
12	
13	
14	

(d) Load Process C

Main memory

0	A.0
1	A.1
2	A.2
3	A.3
4	
5	
6	
7	C.0
8	C.1
9	C.2
10	C.3
11	
12	
13	
14	

(e) Swap out B

Main memory

0	A.0
1	A.1
2	A.2
3	A.3
4	D.0
5	D.1
6	D.2
7	C.0
8	C.1
9	C.2
10	C.3
11	D.3
12	D.4
13	
14	

(f) Load Process D

0	0
1	1
2	2
3	3

Process A
page table

0	N
1	N
2	N

Process B
page table

0	7
1	8
2	9
3	10

Process C
page table

0	4
1	5
2	6
3	11
4	12

Process D

13
14

Free frame
list

页表

Address Translation Scheme

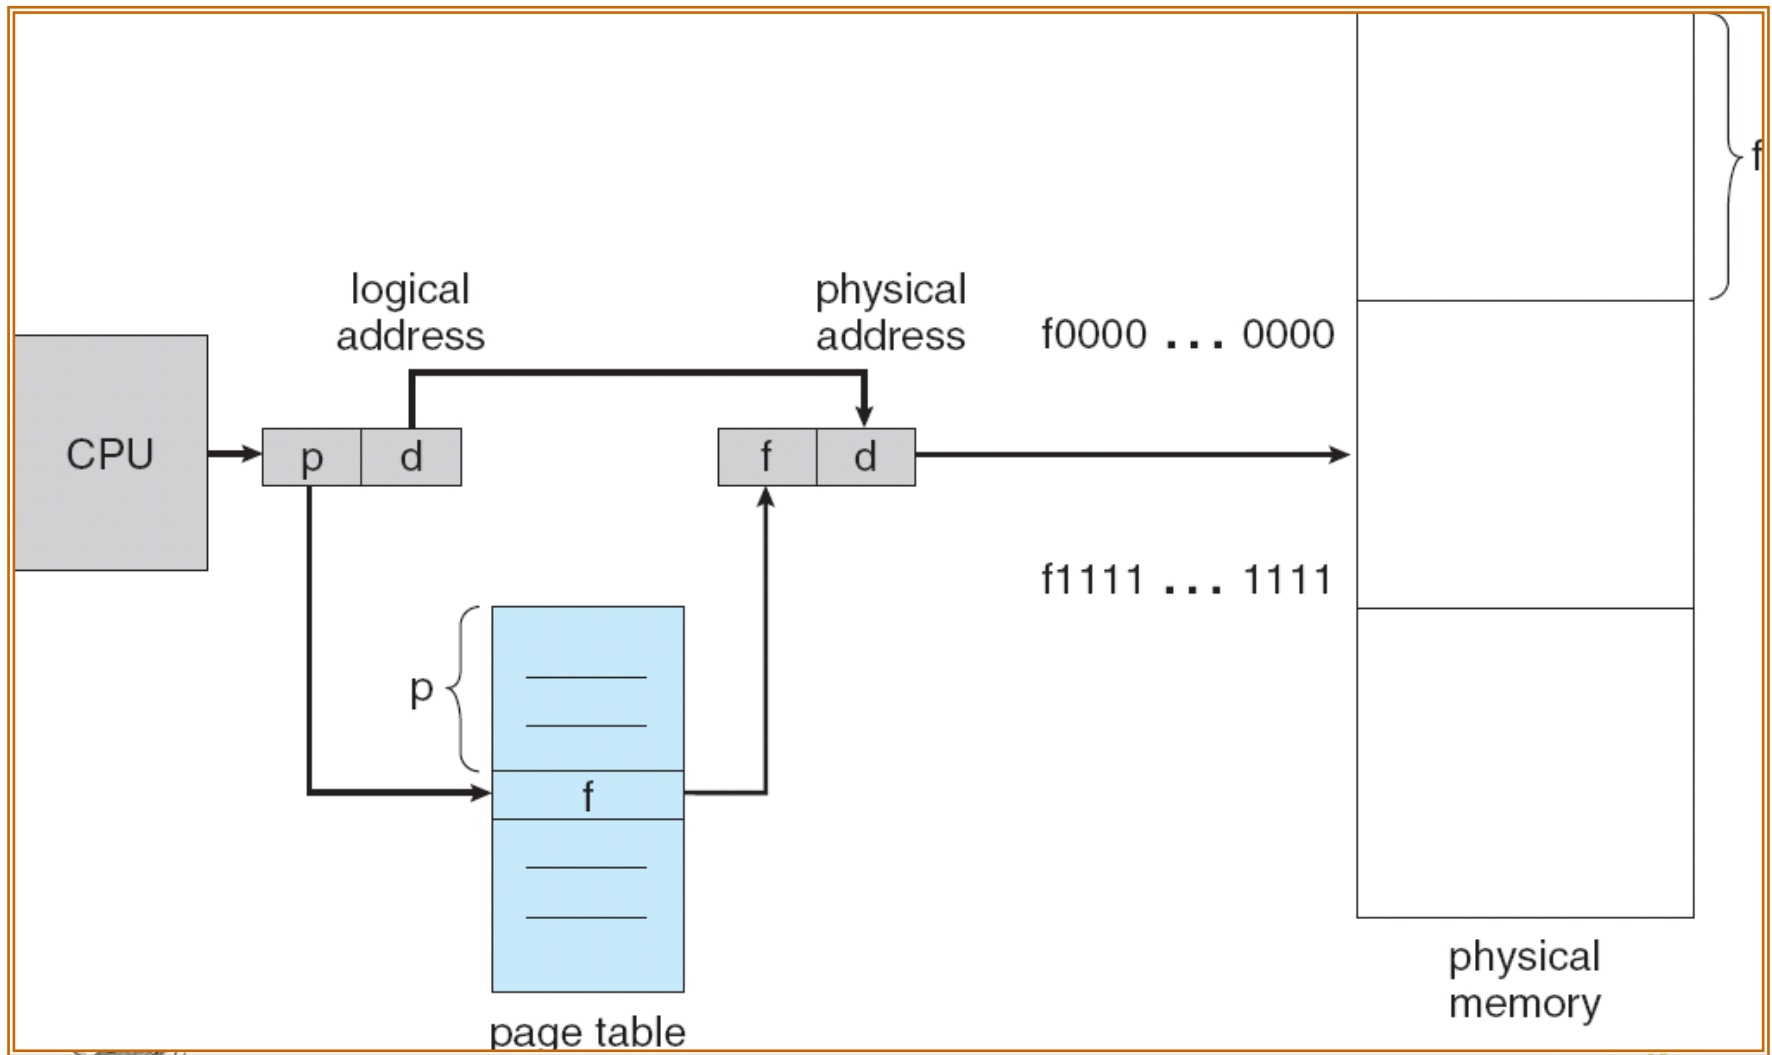
- Address generated by CPU is divided into:
 - **Page number (p)** – used as an index into a *page table* which **contains base address of each page** in physical memory
 - **Page offset (d)** – combined with base address to define the physical memory address that is sent to the memory unit

page number	page offset
p	d
$m - n$	n

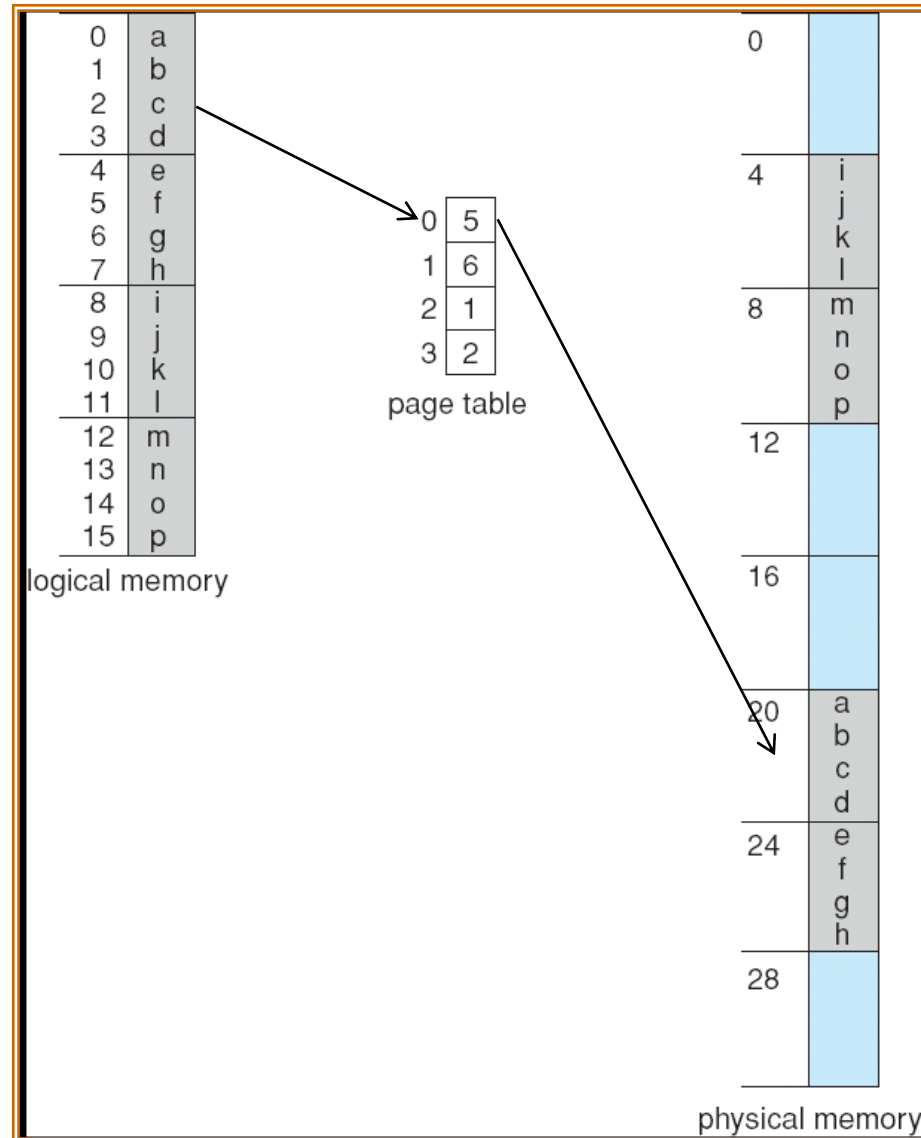
- For given logical address **space 2^m and page size 2^n**



Paging Hardware



Paging Example

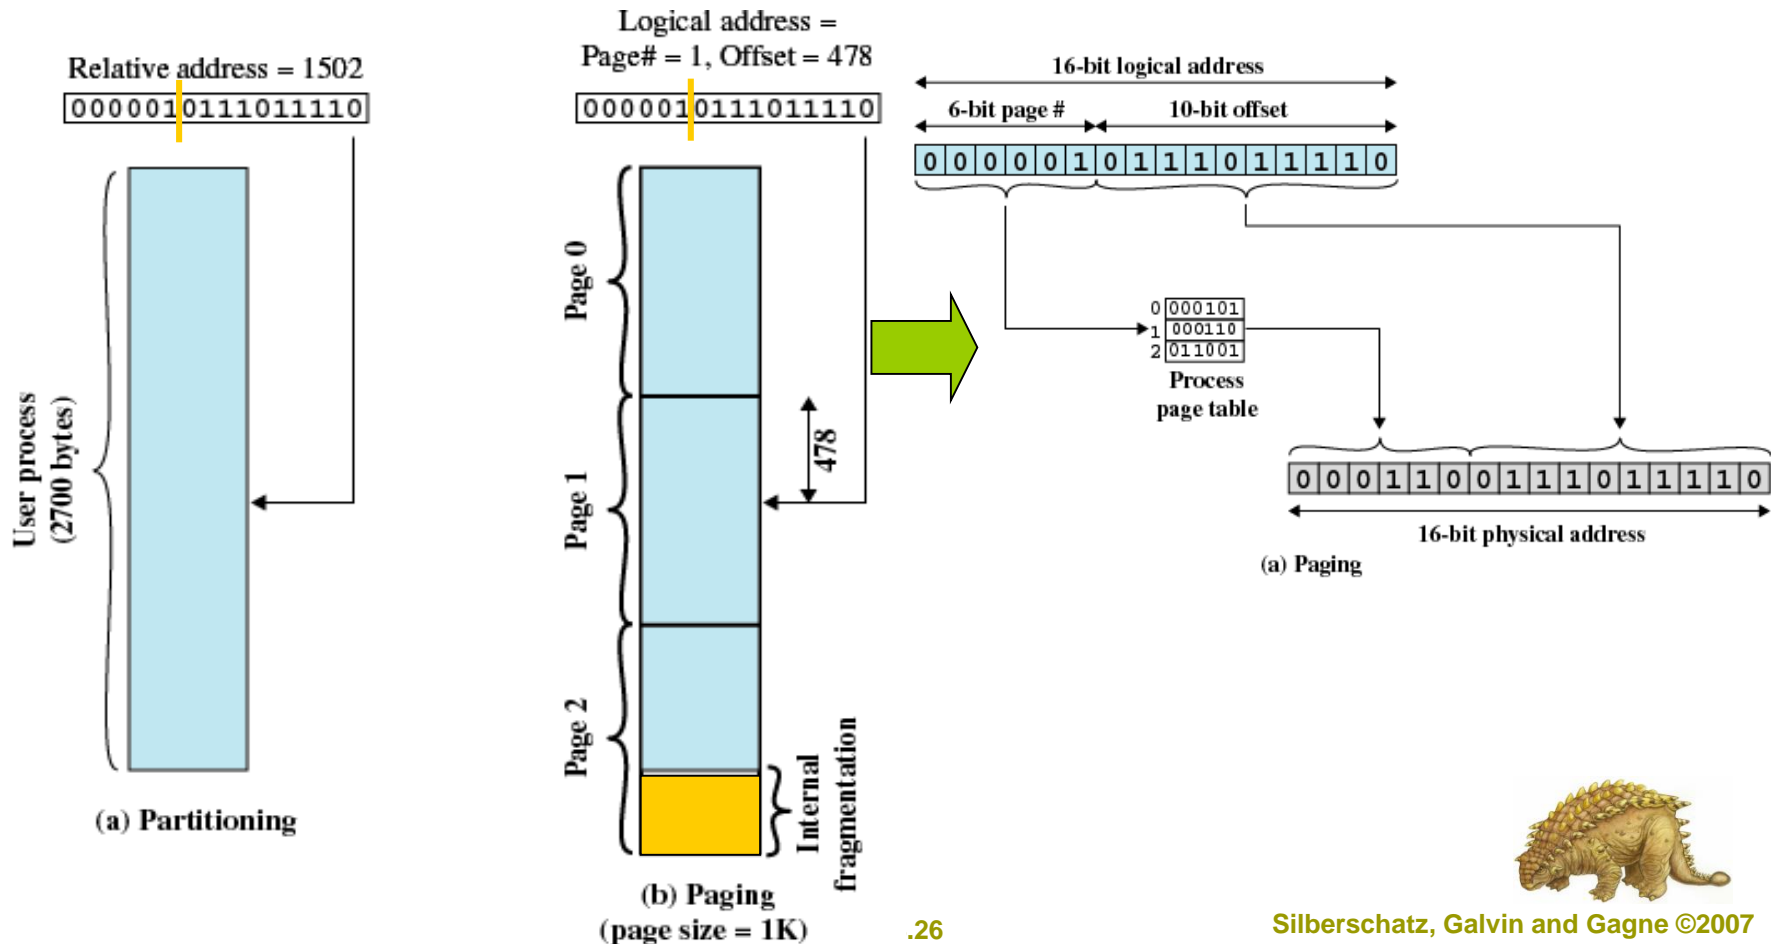


32-byte memory and 4-byte pages



分页技术

- 为了方便使用，规定**页的大小**及**帧的大小**必须是2的幂，以便容易表示出相对地址；
- 相对地址由程序的起点和逻辑地址定义，可以用页号和偏移量表示。



Implementation of Page Table

- ❑ Page table is **kept in main memory**
- ❑ **Page-table base register (PTBR)** **points** to the page table
- ❑ **Page-table length register (PRLR)** indicates **size** of the page table
- ❑ In this scheme every data/instruction access **requires two memory accesses**. **One for the page table** and **one for the data/instruction**.
- ❑ The two memory access problem can be solved by the use of a **special fast-lookup hardware cache** called **associative memory** or **translation look-aside buffers (TLBs)**
- ❑ Some TLBs store **address-space identifiers (ASIDs)** in each TLB entry – uniquely identifies each process to provide address-space protection for that process



Associative Memory

- Associative memory – parallel search

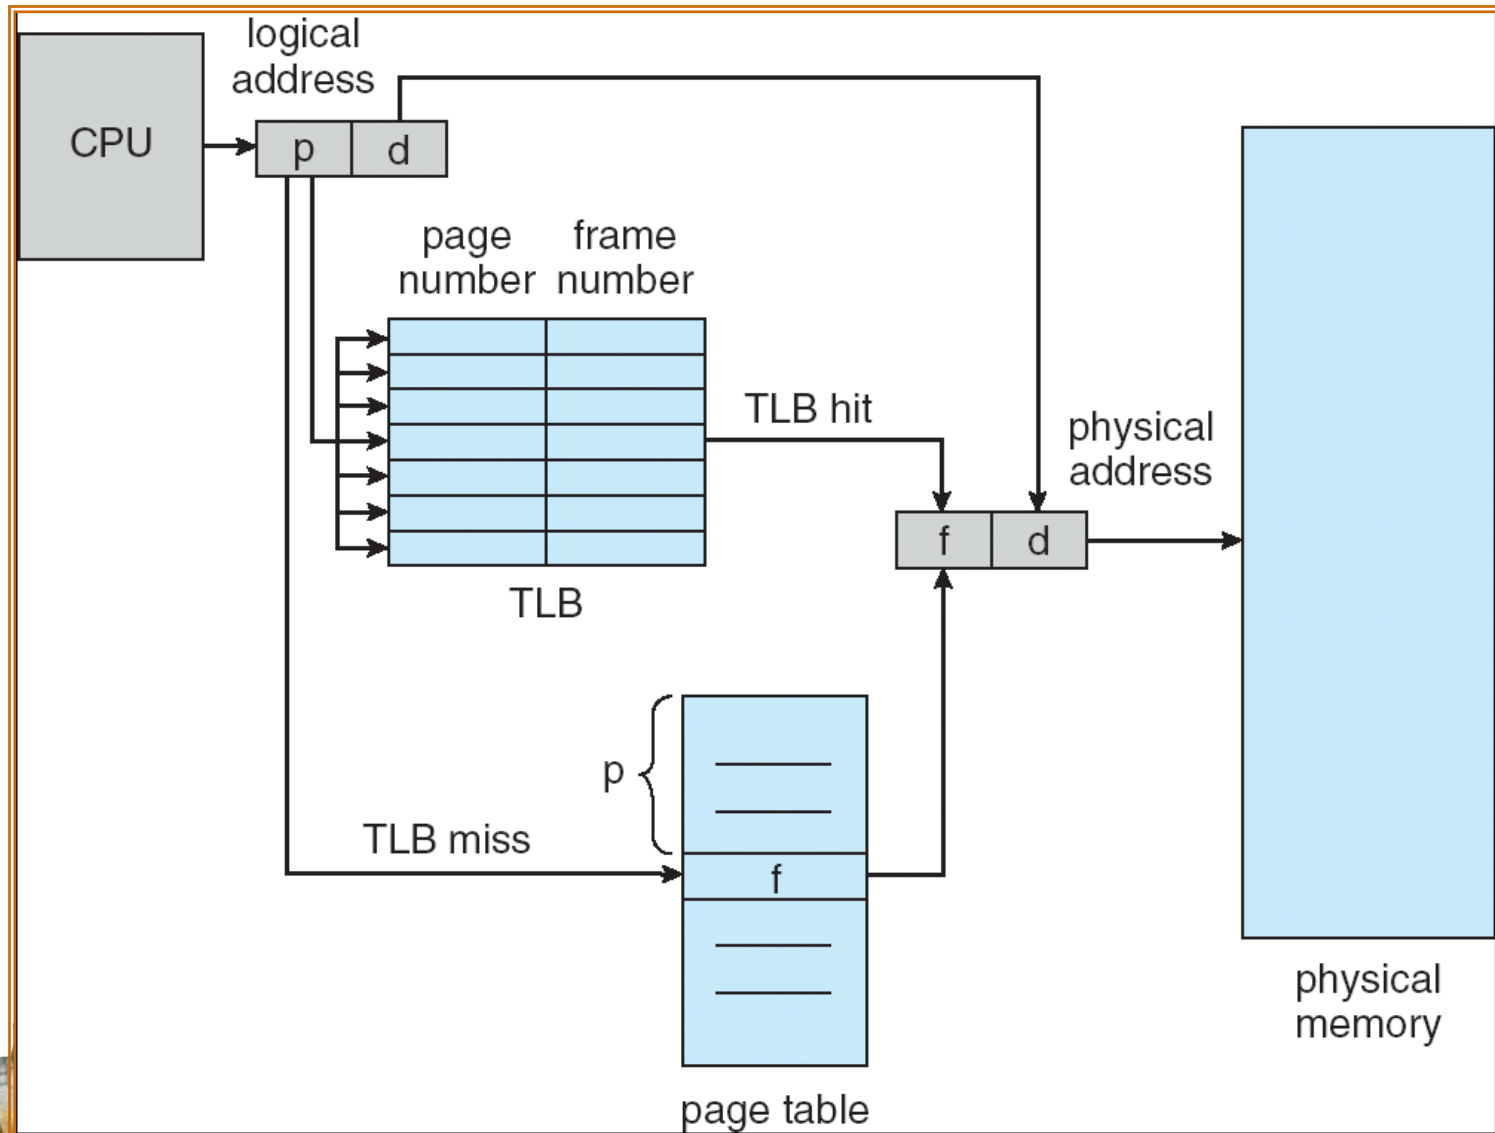
Page #	Frame #

Address translation (p, d)

- If **p** is in associative register, get frame # out
- Otherwise get frame # from page table in memory



Paging Hardware With TLB



Effective Access Time

- ❑ Associative Lookup = ε time unit
- ❑ Assume **memory cycle** time is 1 microsecond
- ❑ **Hit ratio** – percentage of times that a page number **is found** in the associative registers; ratio related to number of associative registers
- ❑ Hit ratio = α
- ❑ **Effective Access Time (EAT)**

$$\begin{aligned} \text{EAT} &= (1 + \varepsilon) \alpha + (2 + \varepsilon)(1 - \alpha) \\ &= 2 + \varepsilon - \alpha \end{aligned}$$

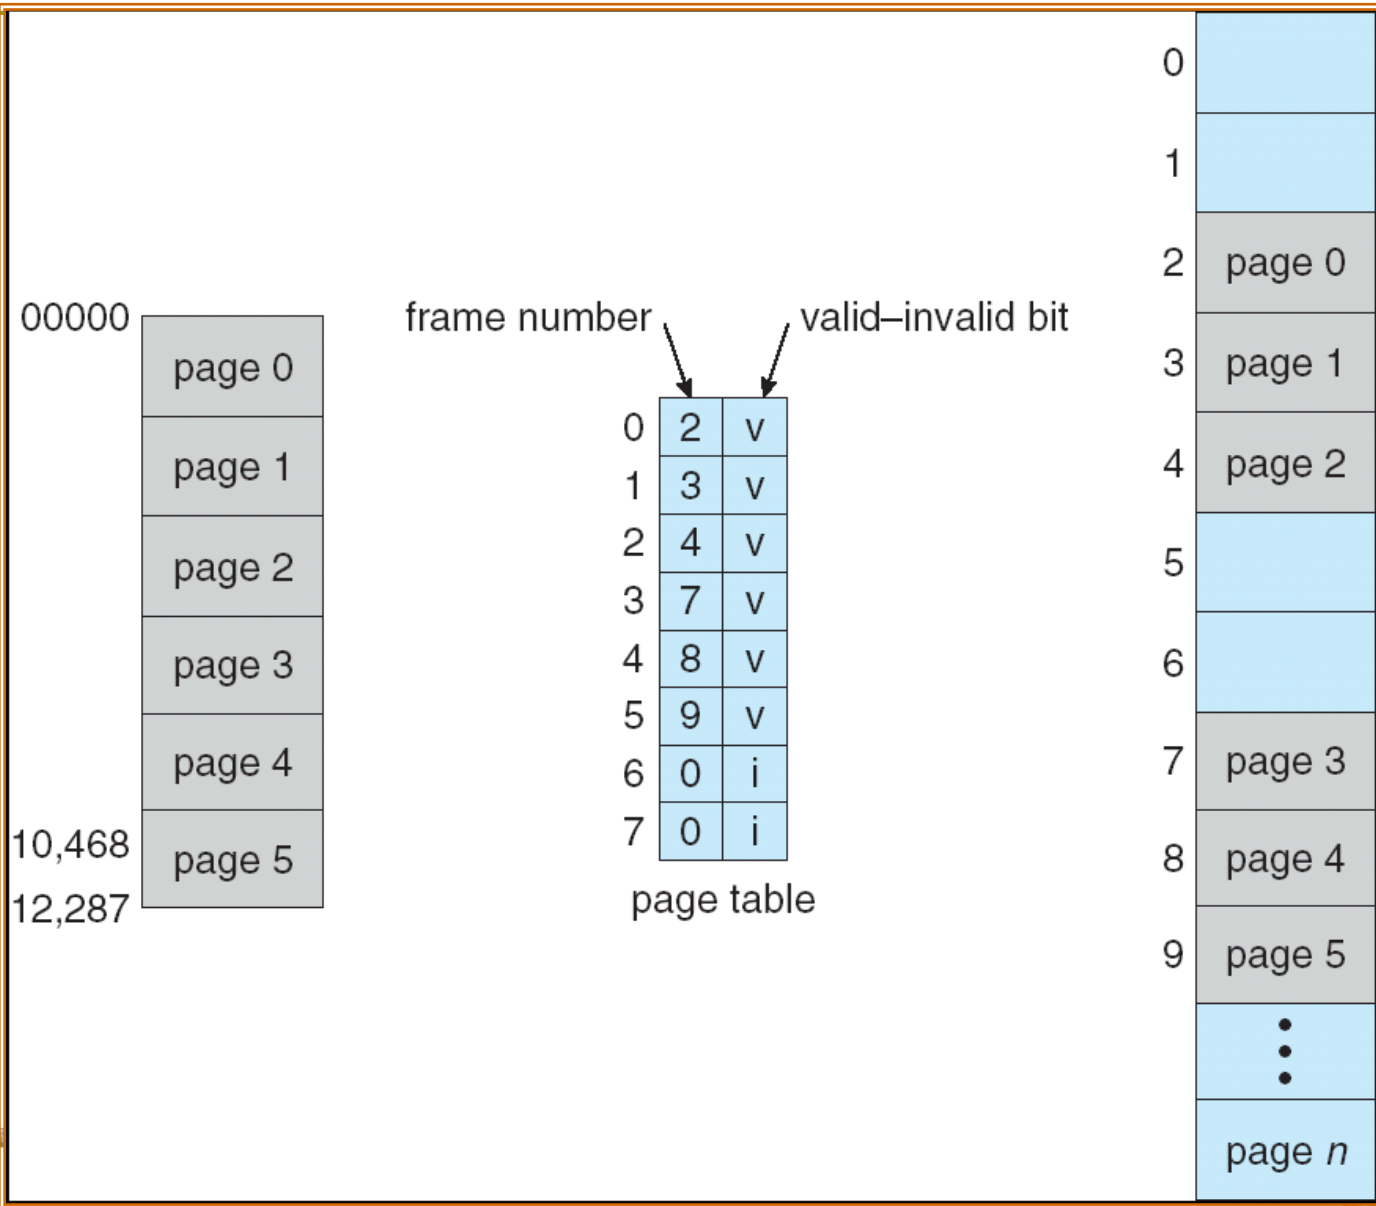


Memory Protection

- ❑ Memory protection implemented by **associating protection bit** with each frame
- ❑ **Valid-invalid** bit attached to each **entry in the page table**:
 - “valid” indicates that **the associated page is in the process’** logical address space, and is thus a legal page
 - “invalid” indicates that the page **is not in the process’ logical address space**



Valid (v) or Invalid (i) Bit In A Page Table



Structure of the Page Table

- ❑ Hierarchical Paging
- ❑ Hashed Page Tables
- ❑ Inverted Page Tables

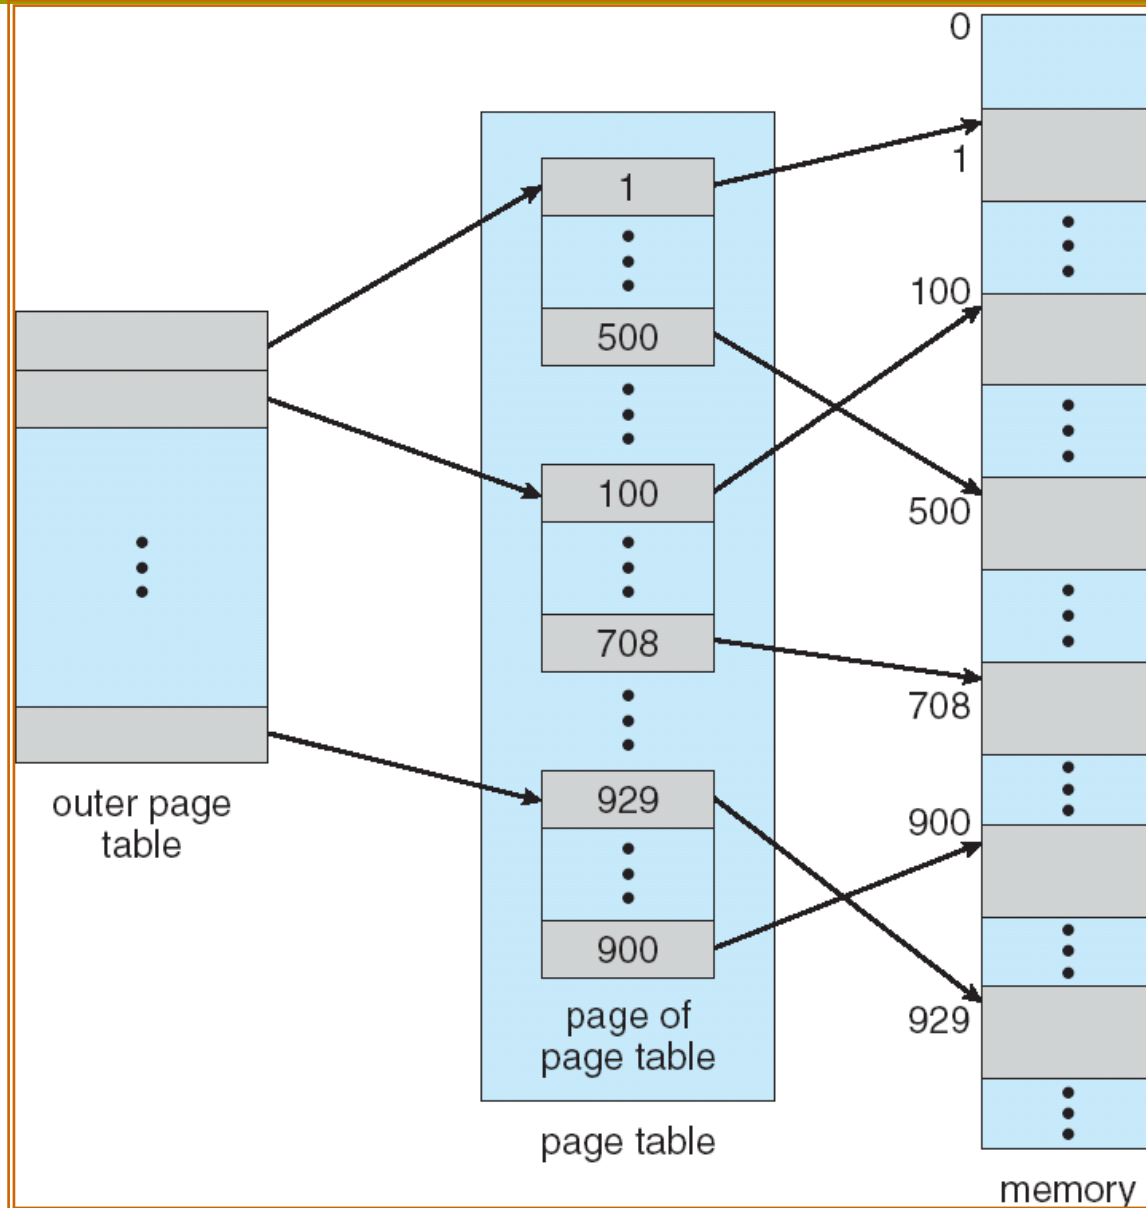


Hierarchical Page Tables

- ❑ Break up the logical address space into **multiple page tables**
- ❑ A simple technique is **a two-level page table**

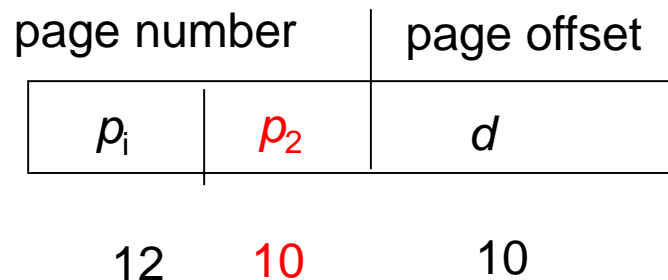


Two-Level Page-Table Scheme



Two-Level Paging Example

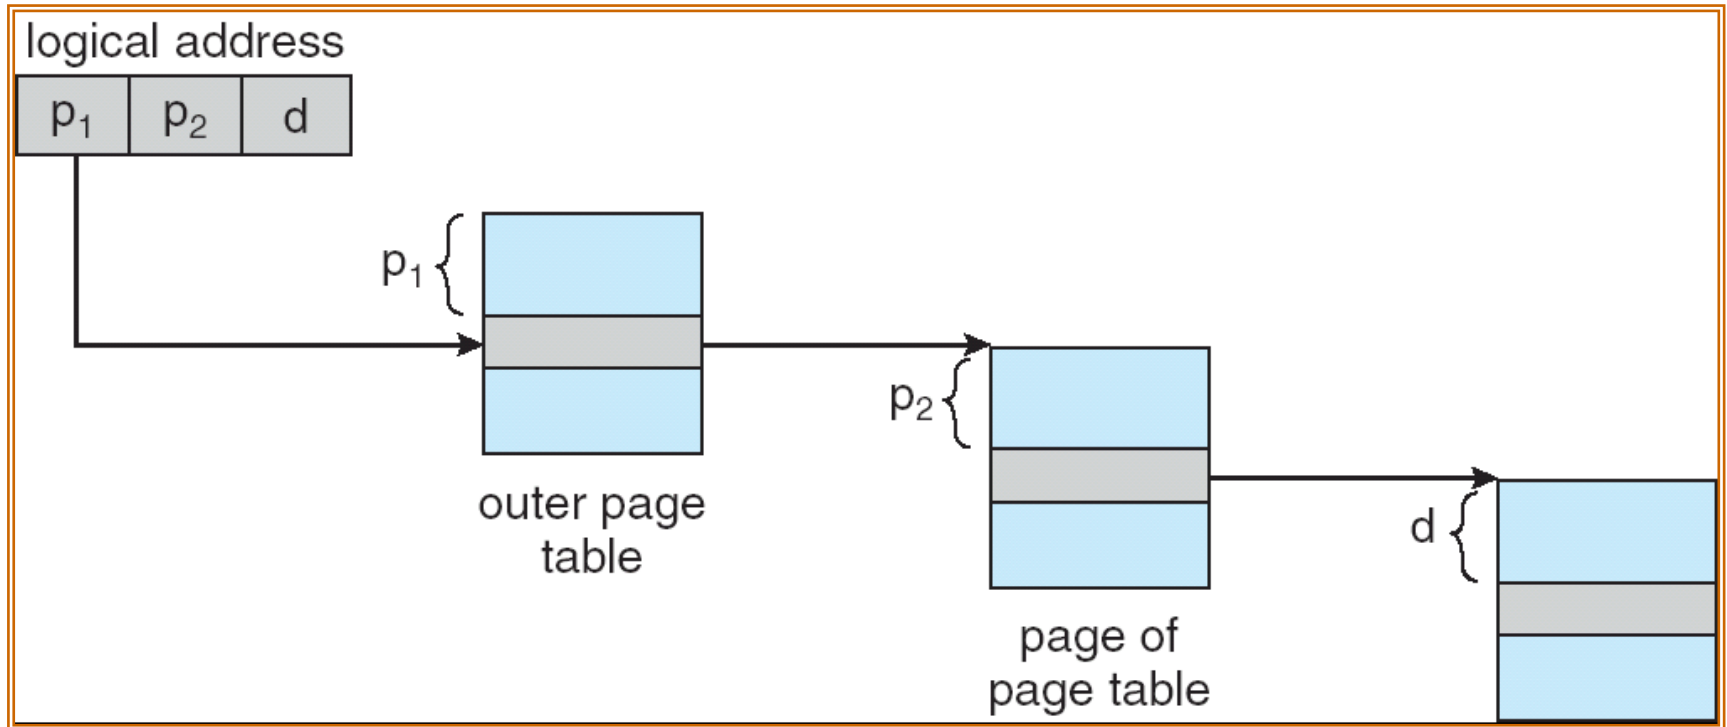
- A logical address (on 32-bit machine with 1K page size) is divided into:
 - a page number consisting of **22 bits**
 - a page offset consisting of **10 bits**
- Since the page table is paged, the page number is further divided into:
 - a **12-bit page number**
 - a **10-bit page offset**
- Thus, a logical address is as follows:



where p_i is an index into the outer page table, and p_2 is the displacement within the page of the outer page table



Address-Translation Scheme



Three-level Paging Scheme

outer page	inner page	offset
p_1	p_2	d
42	10	12

2nd outer page	outer page	inner page	offset
p_1	p_2	p_3	d
32	10	10	12

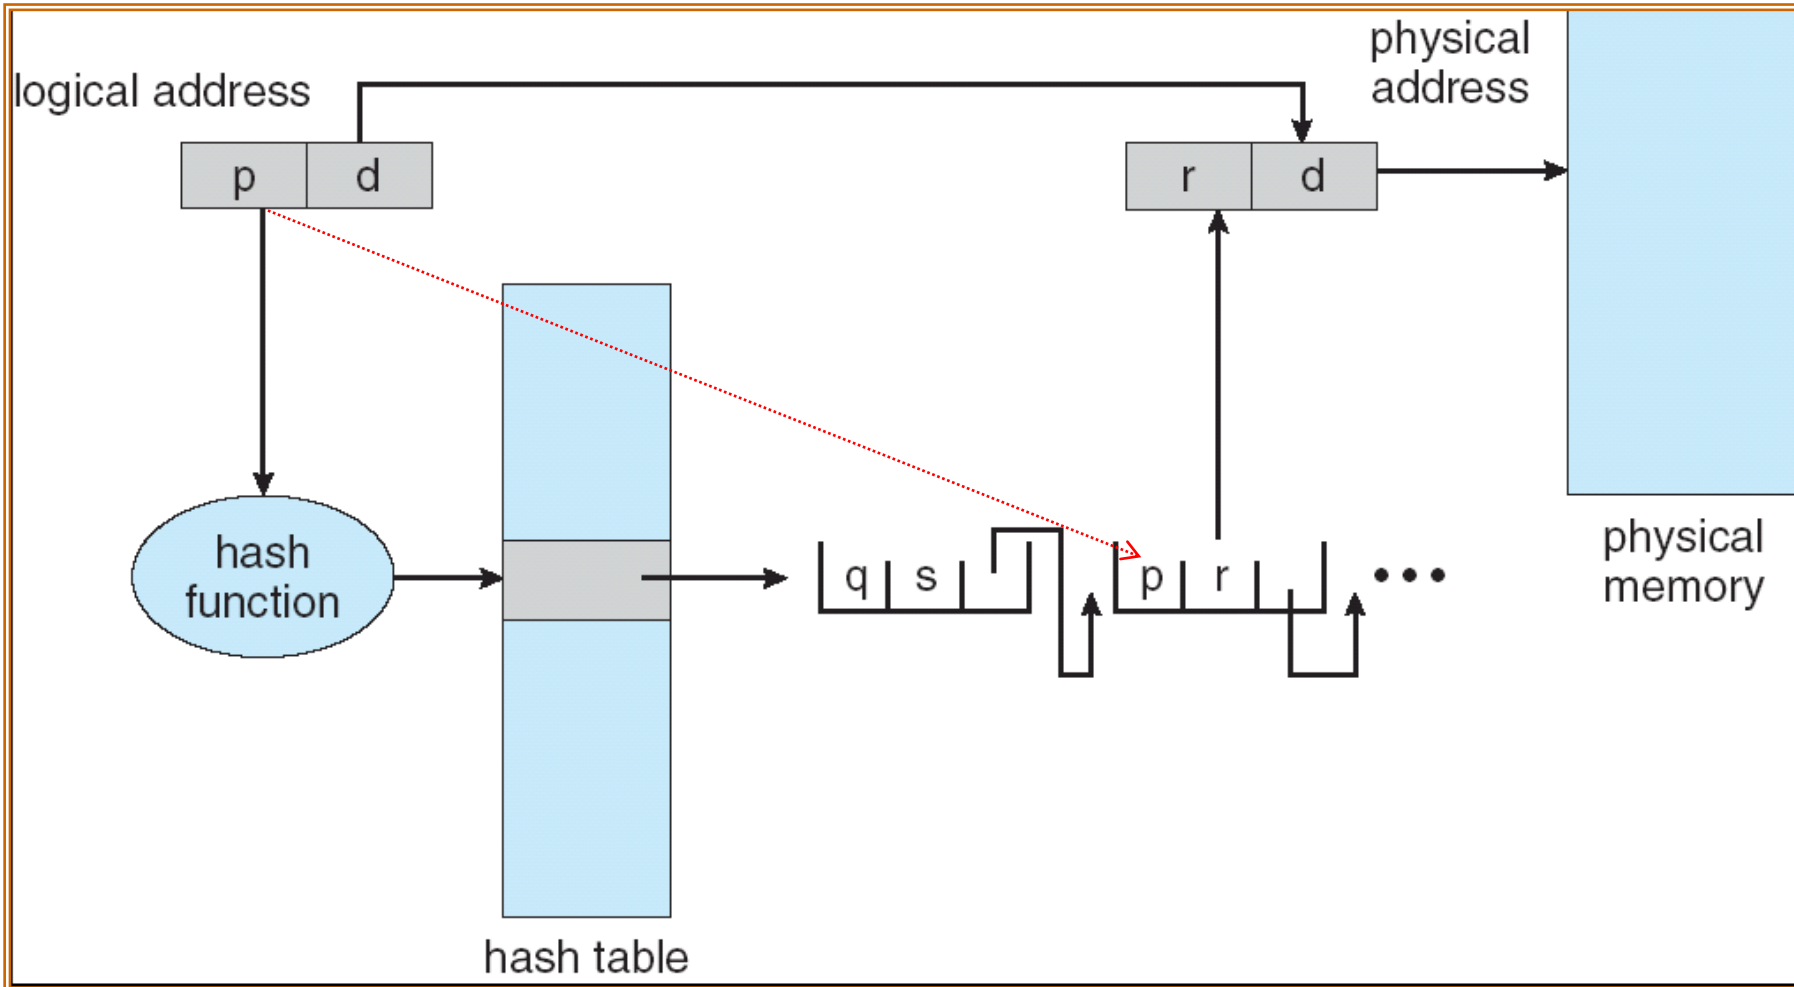


Hashed Page Tables

- ❑ Common in address spaces **> 32 bits**
- ❑ The virtual page number is hashed into a page table. This page table contains a chain of elements hashing to the same location.
- ❑ Virtual page numbers are compared in this chain searching for a match. If a match is found, the corresponding physical frame is extracted.



Hashed Page Table

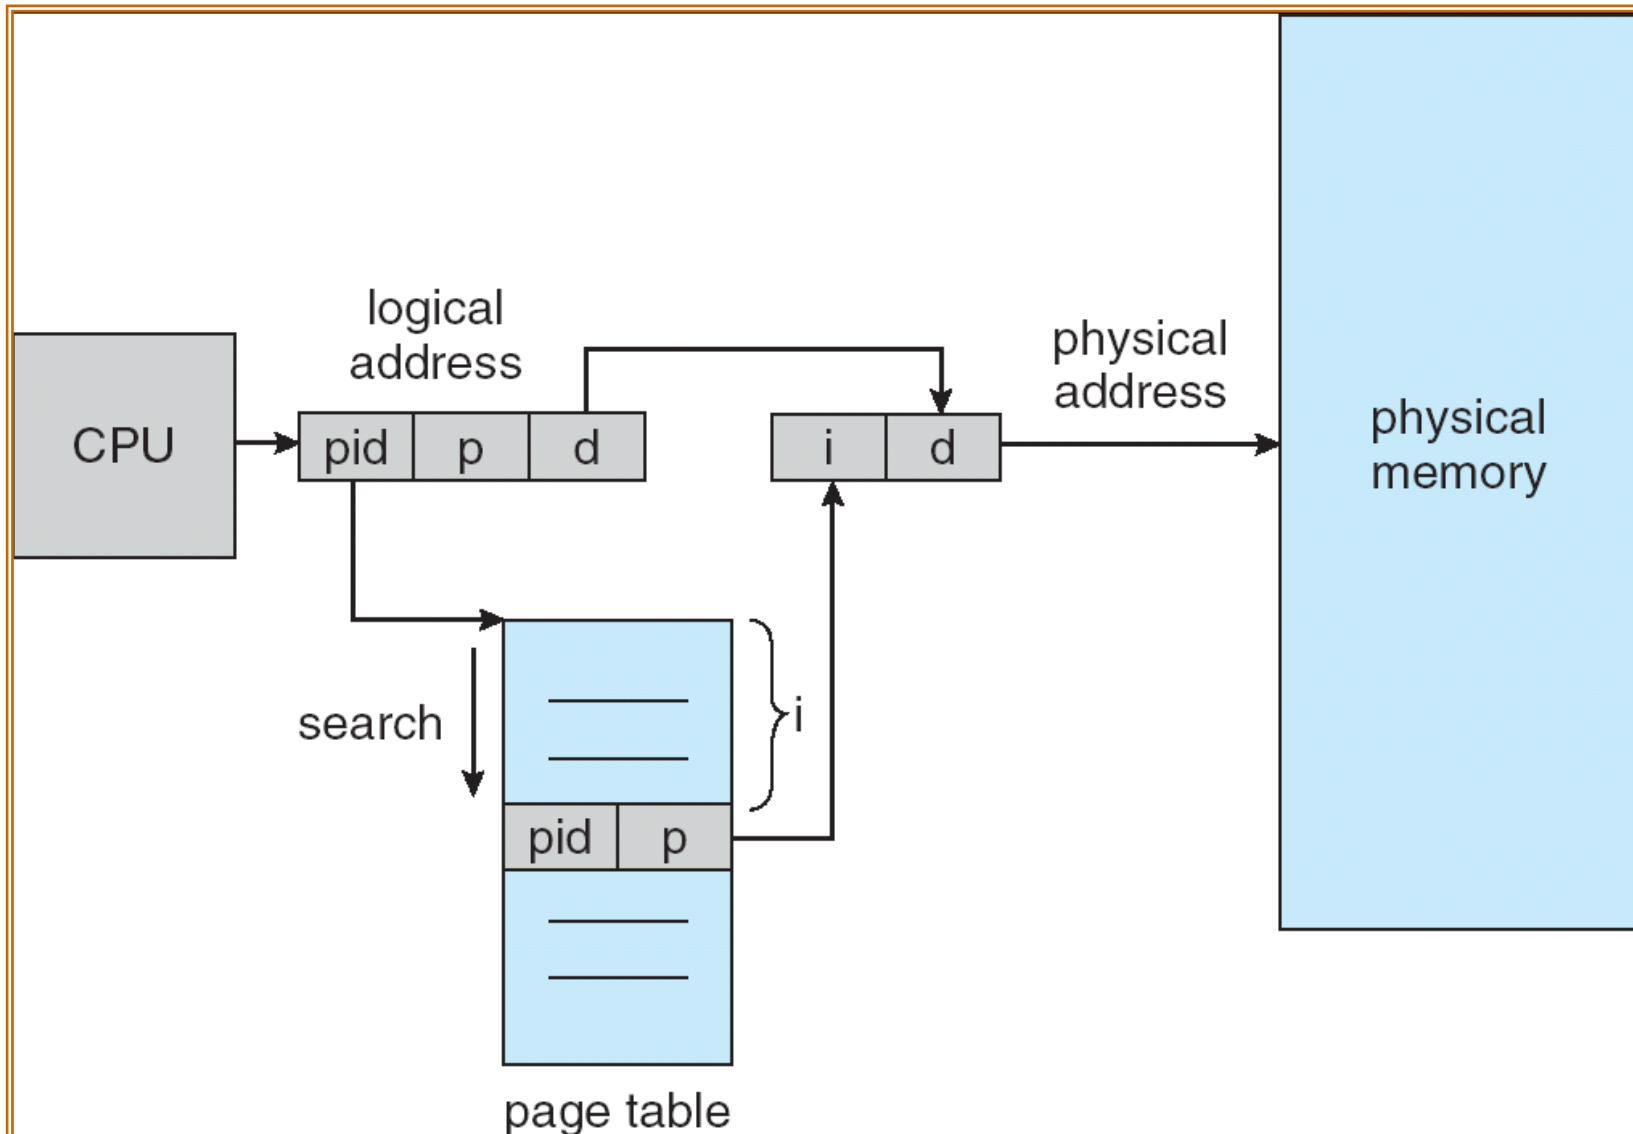


Inverted Page Table

- ❑ One entry for **each real page of memory**
- ❑ Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page
- ❑ Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs
- ❑ Use hash table to limit the search to one — or at most a few — page-table entries



Inverted Page Table Architecture



Shared Pages

❑ Shared code

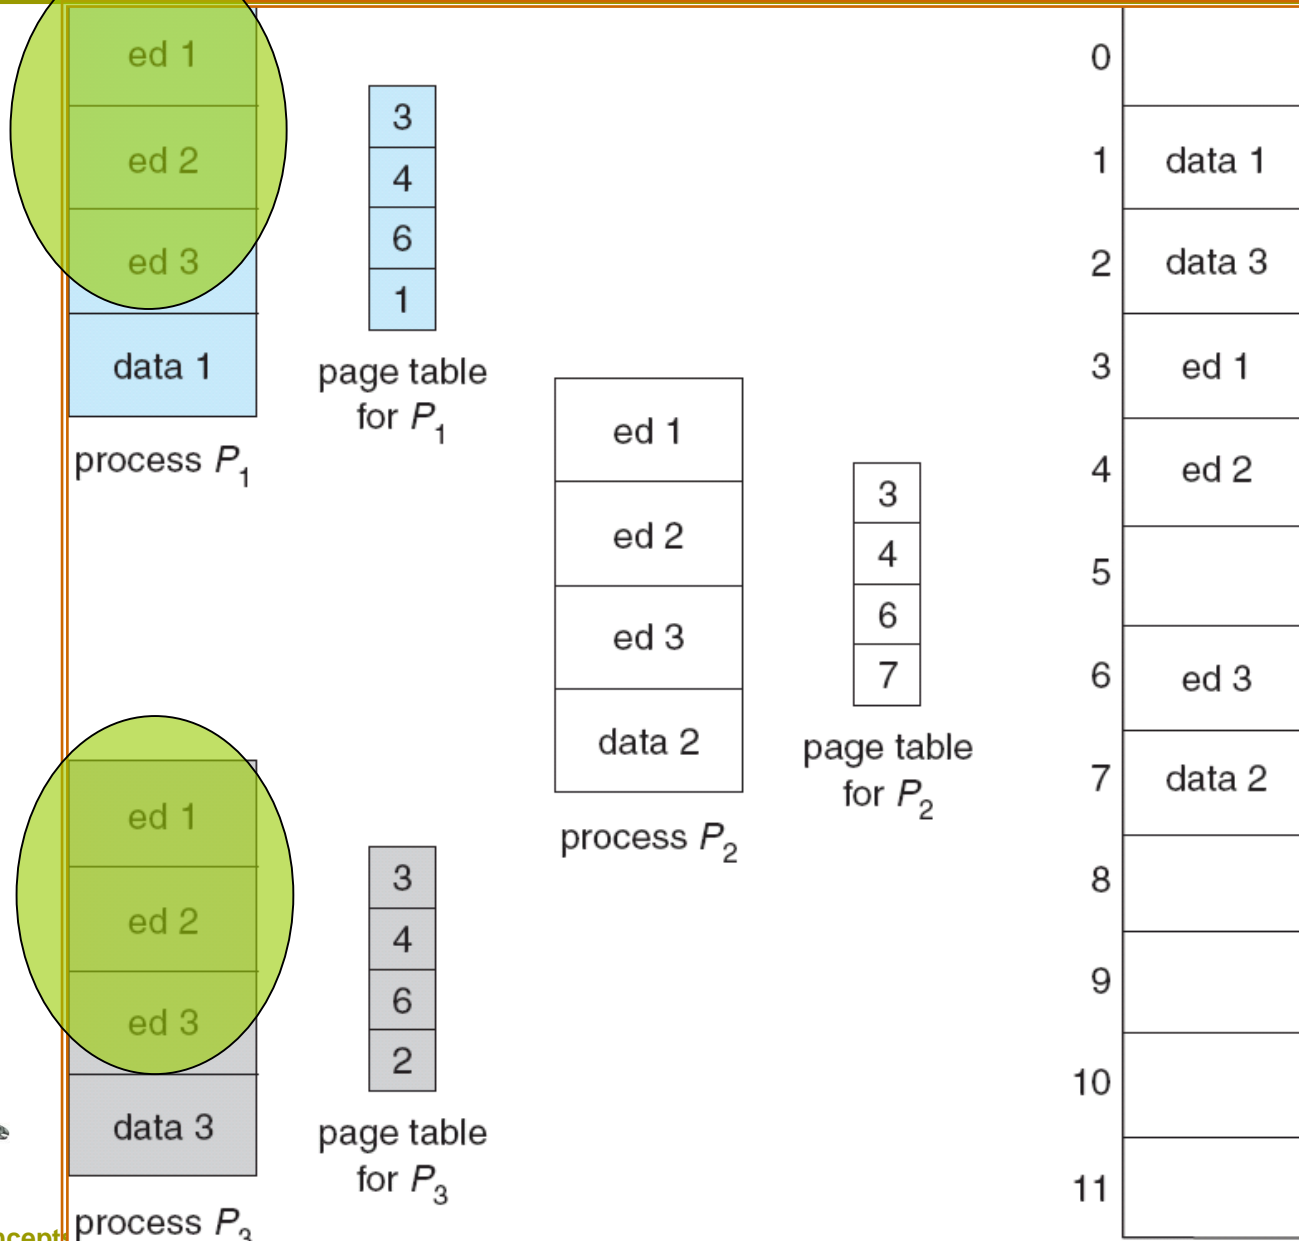
- One copy of read-only (reentrant) code **shared among processes** (i.e., text editors, compilers, window systems).
- Shared code must appear in same location in the logical address space of all processes

❑ Private code and data

- Each process keeps a separate copy of the code and data
- The pages for the private code and data can appear anywhere in the logical address space



Shared Pages Example

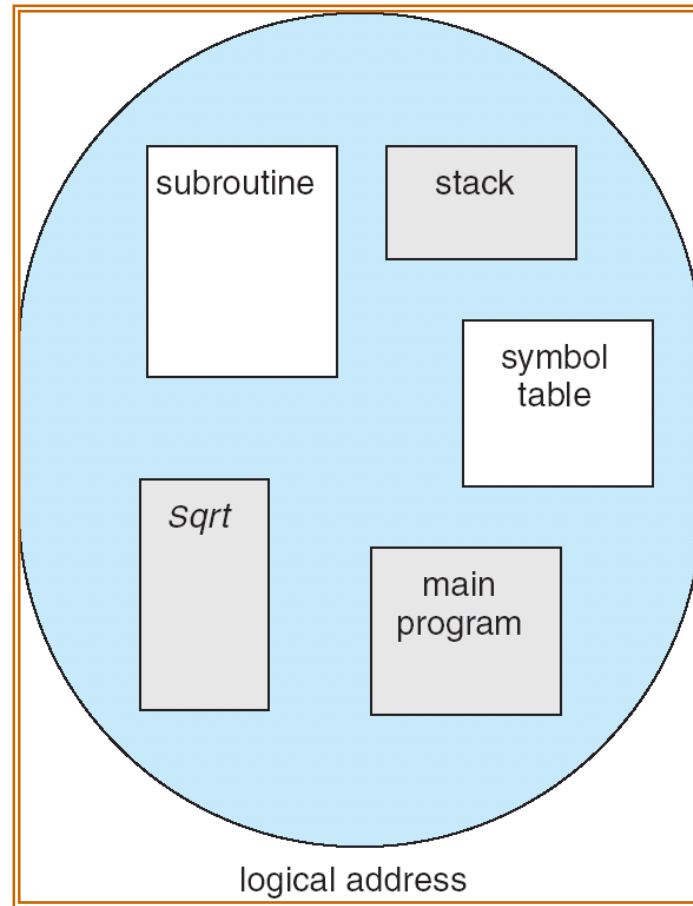


Segmentation

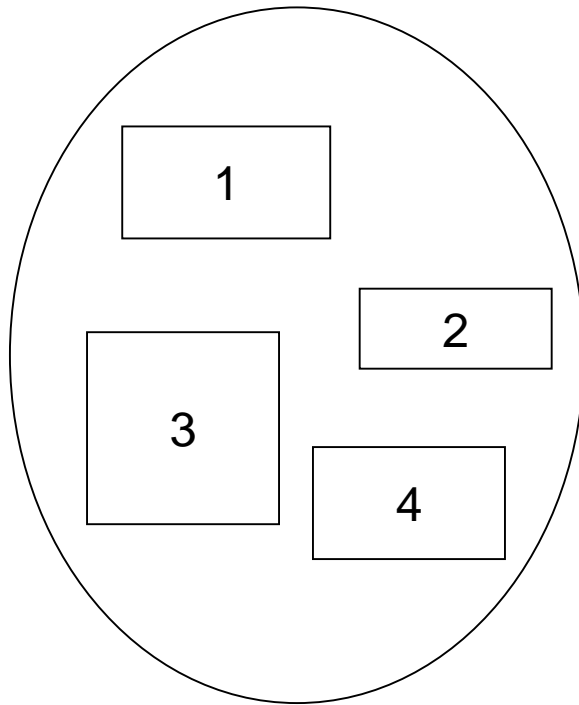
- ❑ Memory-management scheme that supports user view of memory
- ❑ A program is a collection of segments. A segment is a logical unit such as:
 - main program,
 - procedure,
 - function,
 - method,
 - object,
 - local variables, global variables,
 - common block,
 - stack,
 - symbol table, arrays



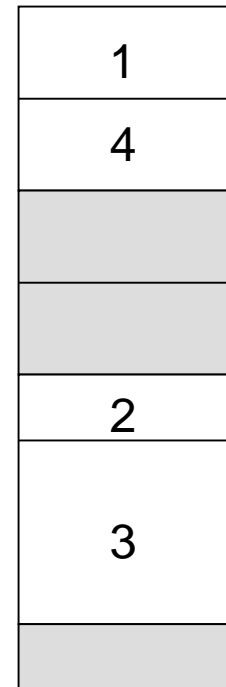
User's View of a Program



Logical View of Segmentation



user space



physical memory space

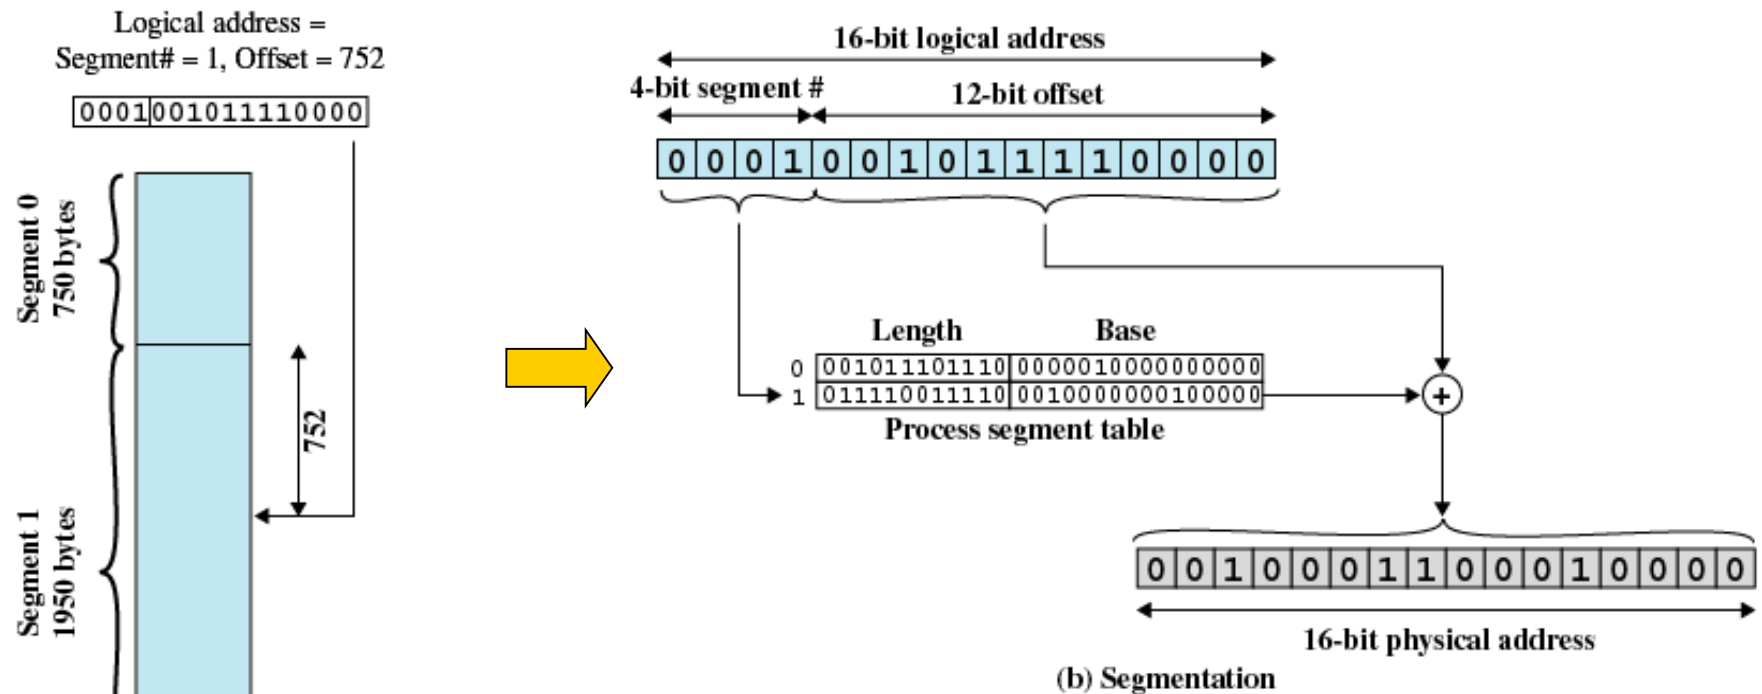


分段技术

- 程序和相关的被**划分成一组段**（segment）。
- 分段技术的逻辑地址由两部分组成：**段号**和**偏移量**。
- 由于使用大小不等的段，分段类似于动态分区。与动态分区不同的是，在分段方案中，**一个程序可以占据多个分区，并且这些分区不要求是连续的**。
- 分段消除了内部碎片，但是和动态分区一样，会产生很多**外部碎片**。
- 由于进程被分为许多小块，**外部碎片较小**。



分段技术



- 一般情况，程序和数据指定到不同的段。
- 上图中，段长度=4，偏移量=12。最大段长度 $2^{12}=4096$
 - 提取段号（最左边4位）；
 - 查找物理地址；
 - 比较偏移量和段长度，是否地址有效；
 - 物理地址=该段起始地址+偏移量



Segmentation Architecture

- Logical address consists of a two tuple:
 <segment-number, offset>,
- **Segment table** – maps two-dimensional physical addresses; each table entry has:
 - **base** – contains the starting physical address where the segments reside in memory
 - **limit** – specifies the length of the segment
- **Segment-table base register (STBR)** points to the segment table's location in memory
- **Segment-table length register (STLR)** indicates number of segments used by a program;
 segment number **s** is legal if **s** < **STLR**

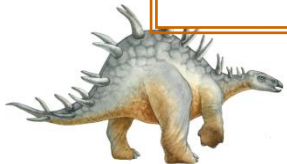
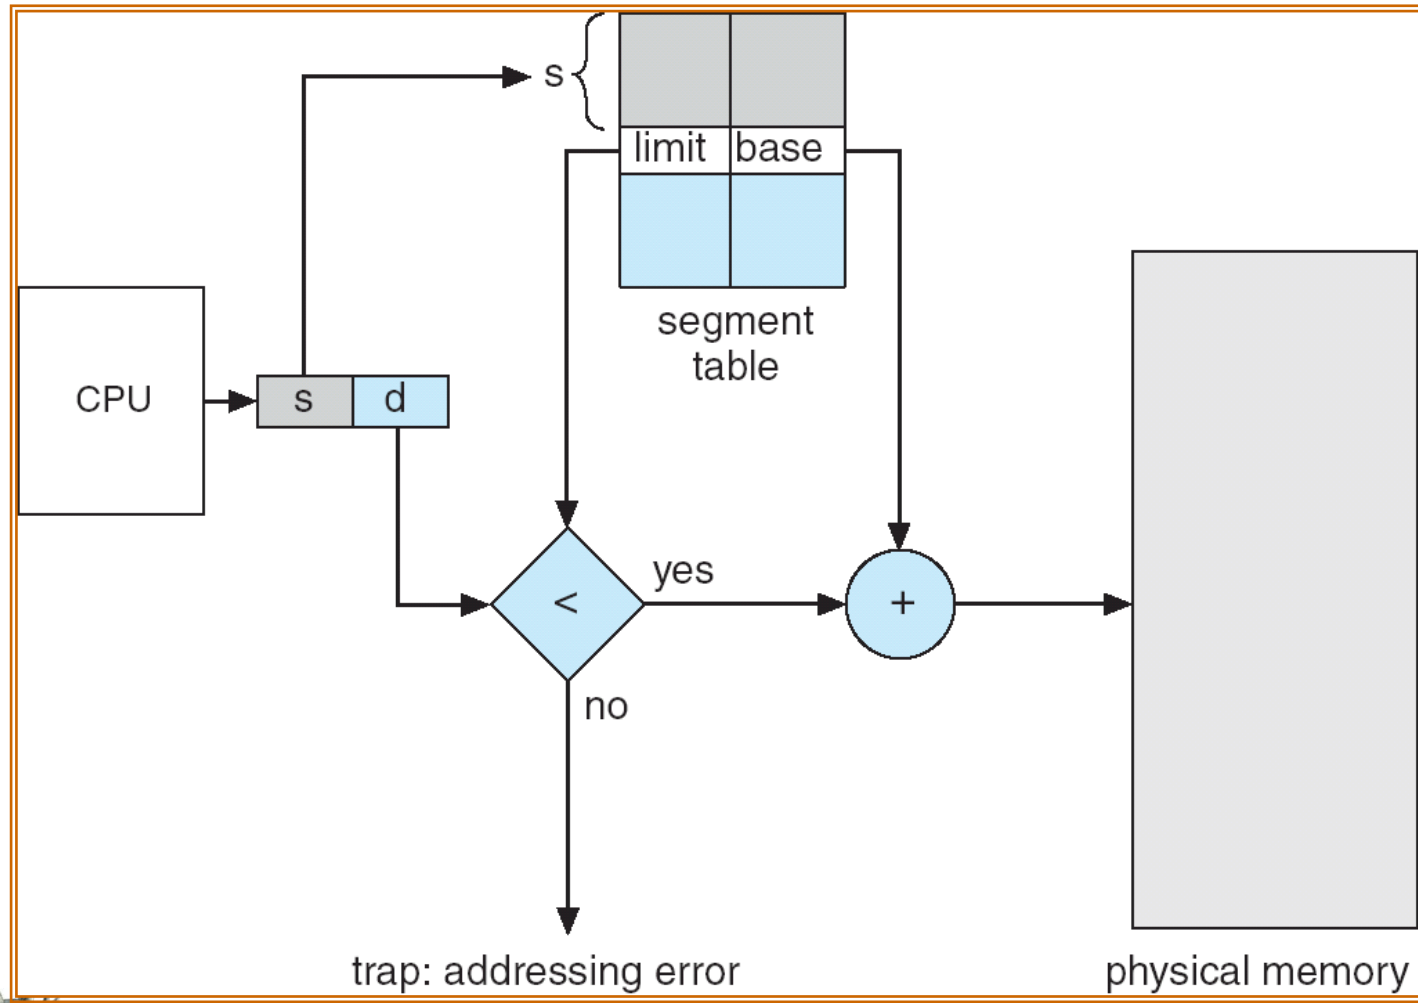


Segmentation Architecture (Cont.)

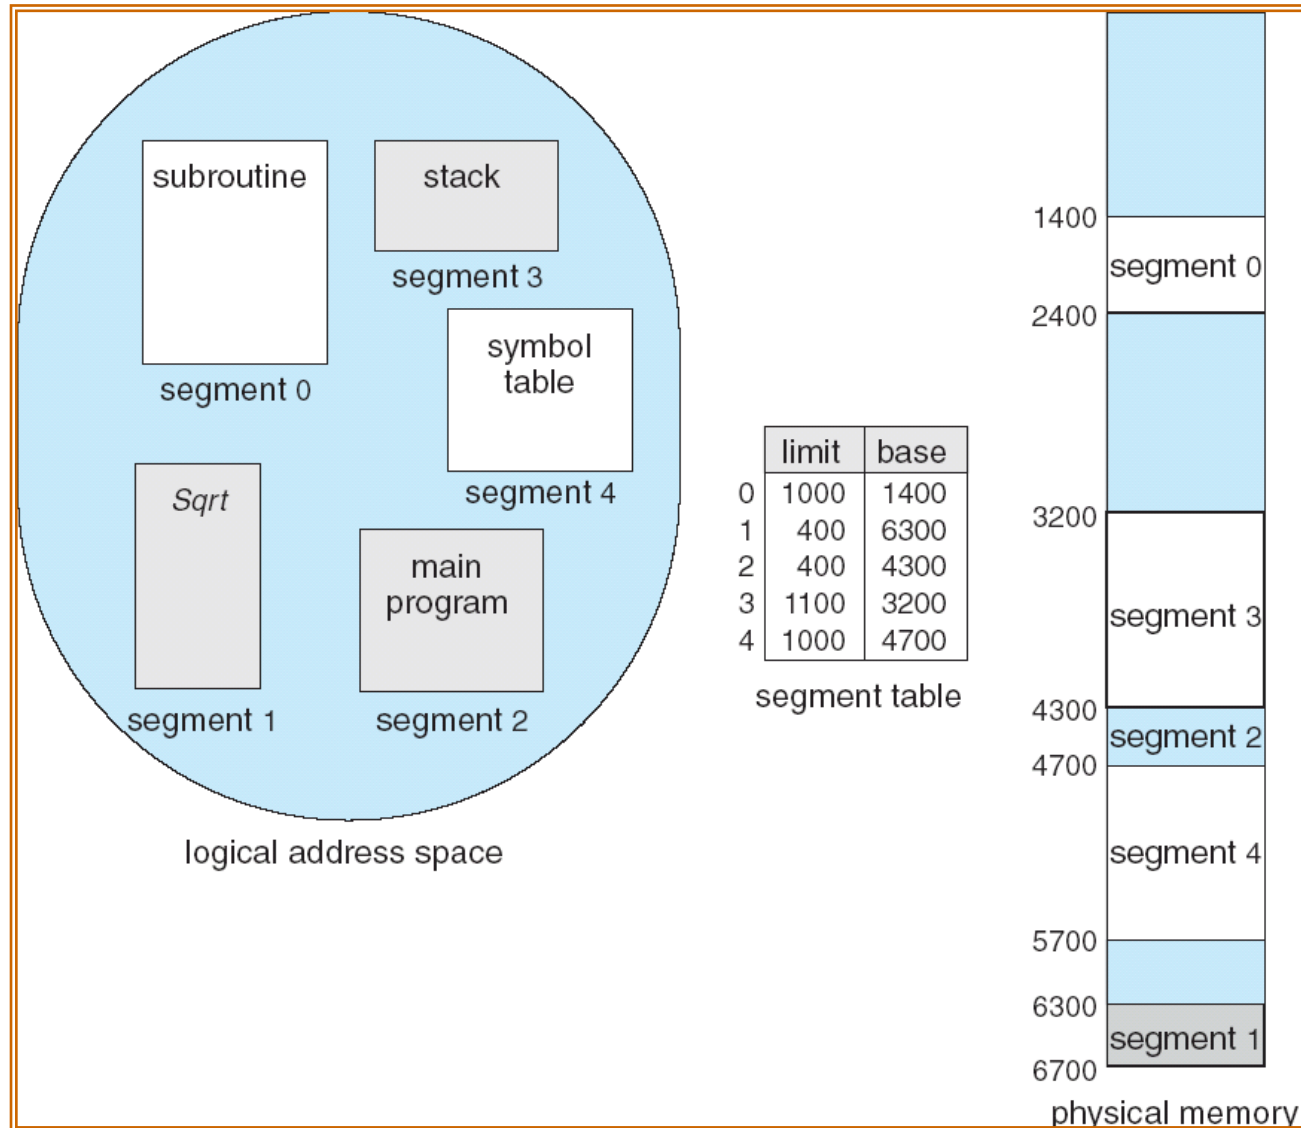
- ❑ Protection
 - With each entry in segment table associate:
 - ❑ validation bit = 0 \Rightarrow **illegal** segment
 - ❑ read/write/execute privileges
- ❑ Protection bits associated with segments; **code sharing occurs at segment level**
- ❑ Since segments vary in length, memory allocation is a dynamic storage-allocation problem
- ❑ A segmentation example is shown in the following diagram



Segmentation Hardware



Example of Segmentation



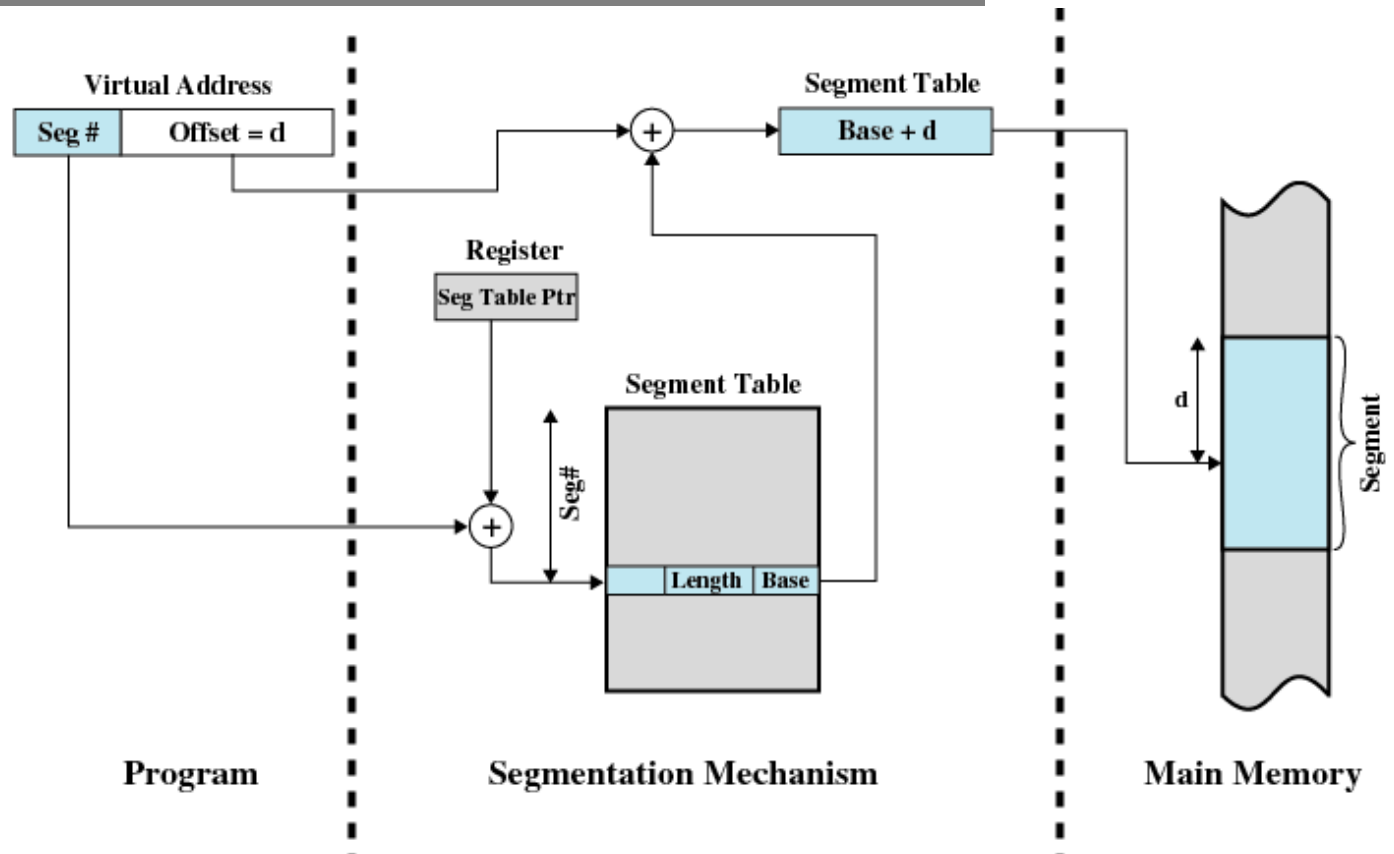
分段式虚拟内存

Segment Number	Offset
----------------	--------

仅仅分段

Segment Table Entry

P	M	Other Control Bits	Length	Segment Base
---	---	--------------------	--------	--------------



分段和分页结合

用户的地址空间被程序员划分为多个段，段由大小相同的若干页构成，页面有页内地址。

段页式管理系统的地址结构为 (S, P, d) ，其中 S 为段号， P 为段内页号， d 为页内偏移量。

Virtual Address



Segment Table Entry



Page Table Entry

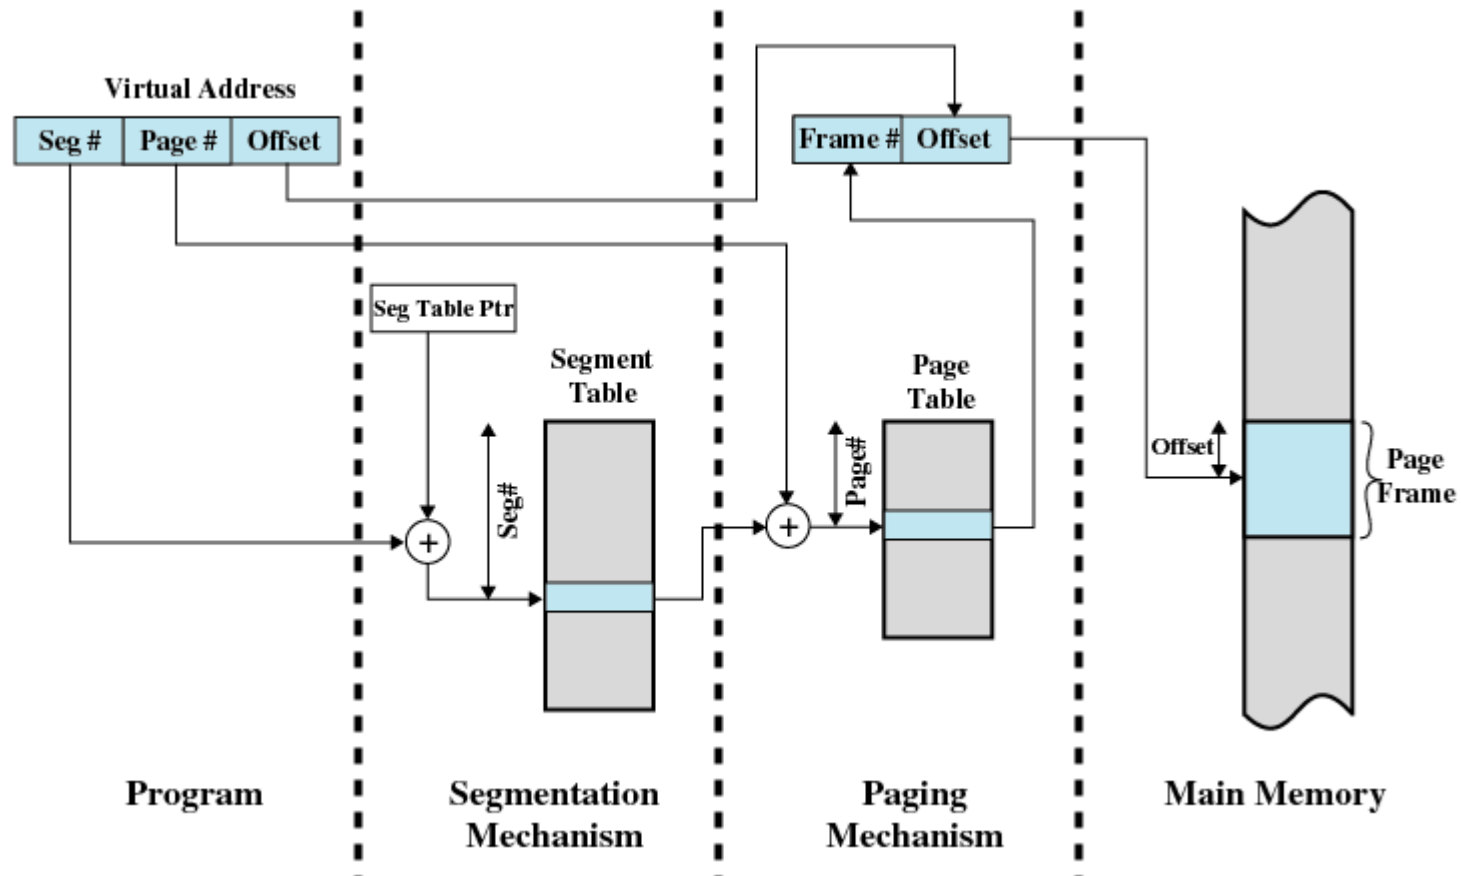


P = present bit
M = Modified bit

(c) Combined segmentation and paging



分段和分页结合



End of Chapter 8

