

Chapter 13: I/O Systems



Chapter 13: I/O Systems

- ❑ I/O Hardware
- ❑ Application I/O Interface
- ❑ Kernel I/O Subsystem
- ❑ Transforming I/O Requests to Hardware Operations
- ❑ Streams
- ❑ Performance



Objectives

- ❑ Explore the structure of an operating system's I/O subsystem
- ❑ Discuss the principles of I/O hardware and its complexity
- ❑ Provide details of the performance aspects of I/O hardware and software




I/O Hardware

- ❑ Incredible variety of I/O devices
- ❑ Common concepts
 - **Port**
 - **Bus** (**daisy chain** or shared direct access)
 - **Controller** (**host adapter**)
- ❑ I/O instructions control devices
- ❑ **Devices have addresses**, used by
 - Direct I/O instructions
 - **Memory-mapped I/O**



A Typical PC Bus Structure



Device I/O Port Locations on PCs (partial)

I/O address range (hexadecimal)	device
000–00F	DMA controller
020–021	interrupt controller
040–043	timer
200–20F	game controller
2F8–2FF	serial port (secondary)
320–32F	hard-disk controller
378–37F	parallel port
3D0–3DF	graphics controller
3F0–3F7	diskette-drive controller
3F8–3FF	serial port (primary)



Polling

- ❑ Determines **state of device**
 - command-ready
 - busy
 - Error
- ❑ **Busy-wait** cycle to wait for I/O from device

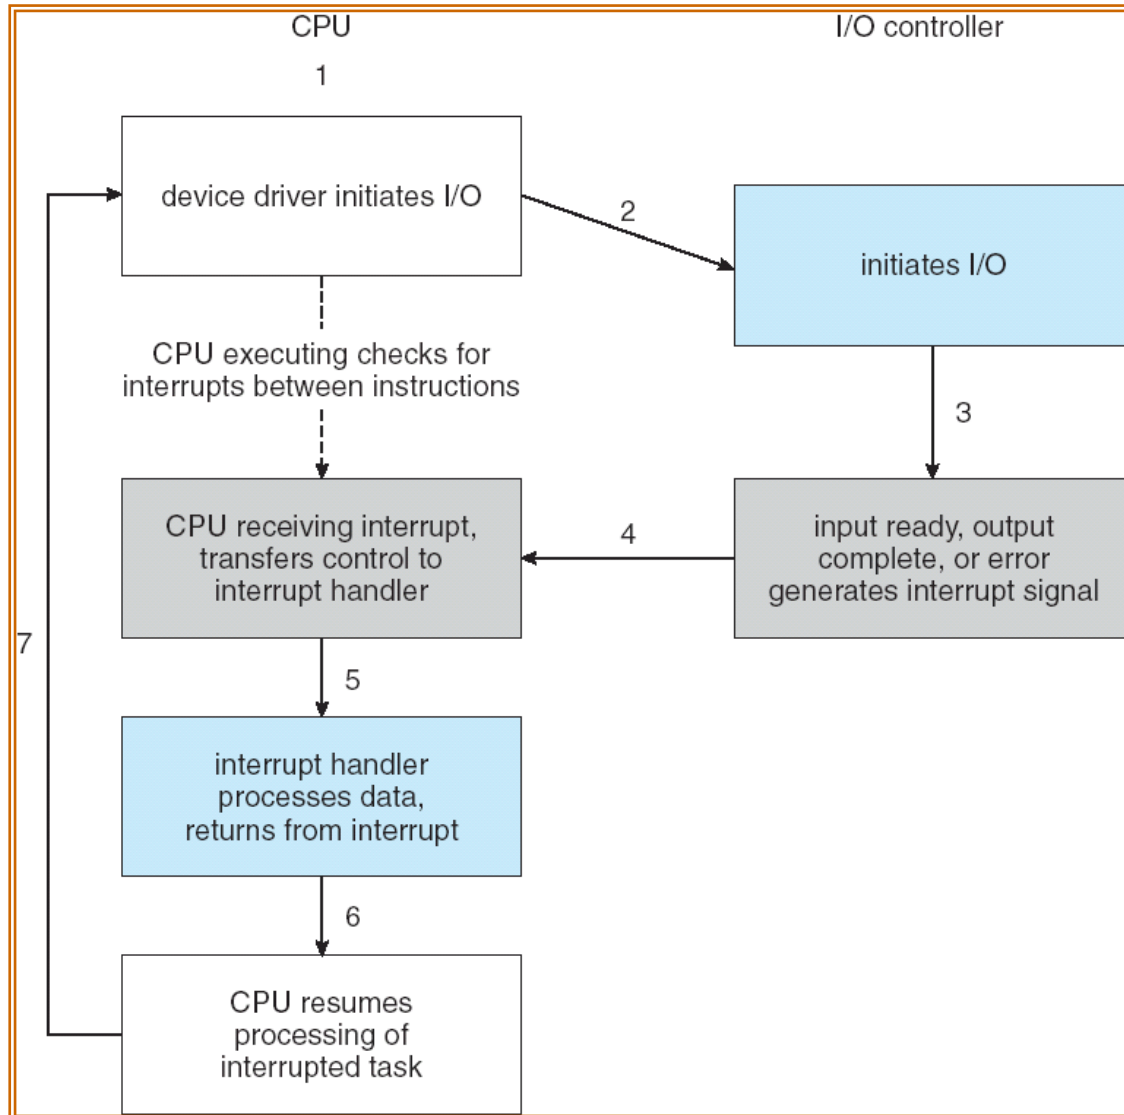


Interrupts

- ❑ CPU **Interrupt-request line** triggered by I/O device
- ❑ **Interrupt handler** receives interrupts
- ❑ **Maskable** to ignore or delay some interrupts
- ❑ Interrupt vector to dispatch interrupt to correct handler
 - Based on priority
 - Some **nonmaskable**
- ❑ Interrupt mechanism also used for exceptions



Interrupt-Driven I/O Cycle



Intel Pentium Processor Event-Vector Table

vector number	description
0	divide error
1	debug exception
2	null interrupt
3	breakpoint
4	INTO-detected overflow
5	bound range exception
6	invalid opcode
7	device not available
8	double fault
9	coprocessor segment overrun (reserved)
10	invalid task state segment
11	segment not present
12	stack fault
13	general protection
14	page fault
15	(Intel reserved, do not use)
16	floating-point error
17	alignment check
18	machine check
19–31	(Intel reserved, do not use)
32–255	maskable interrupts

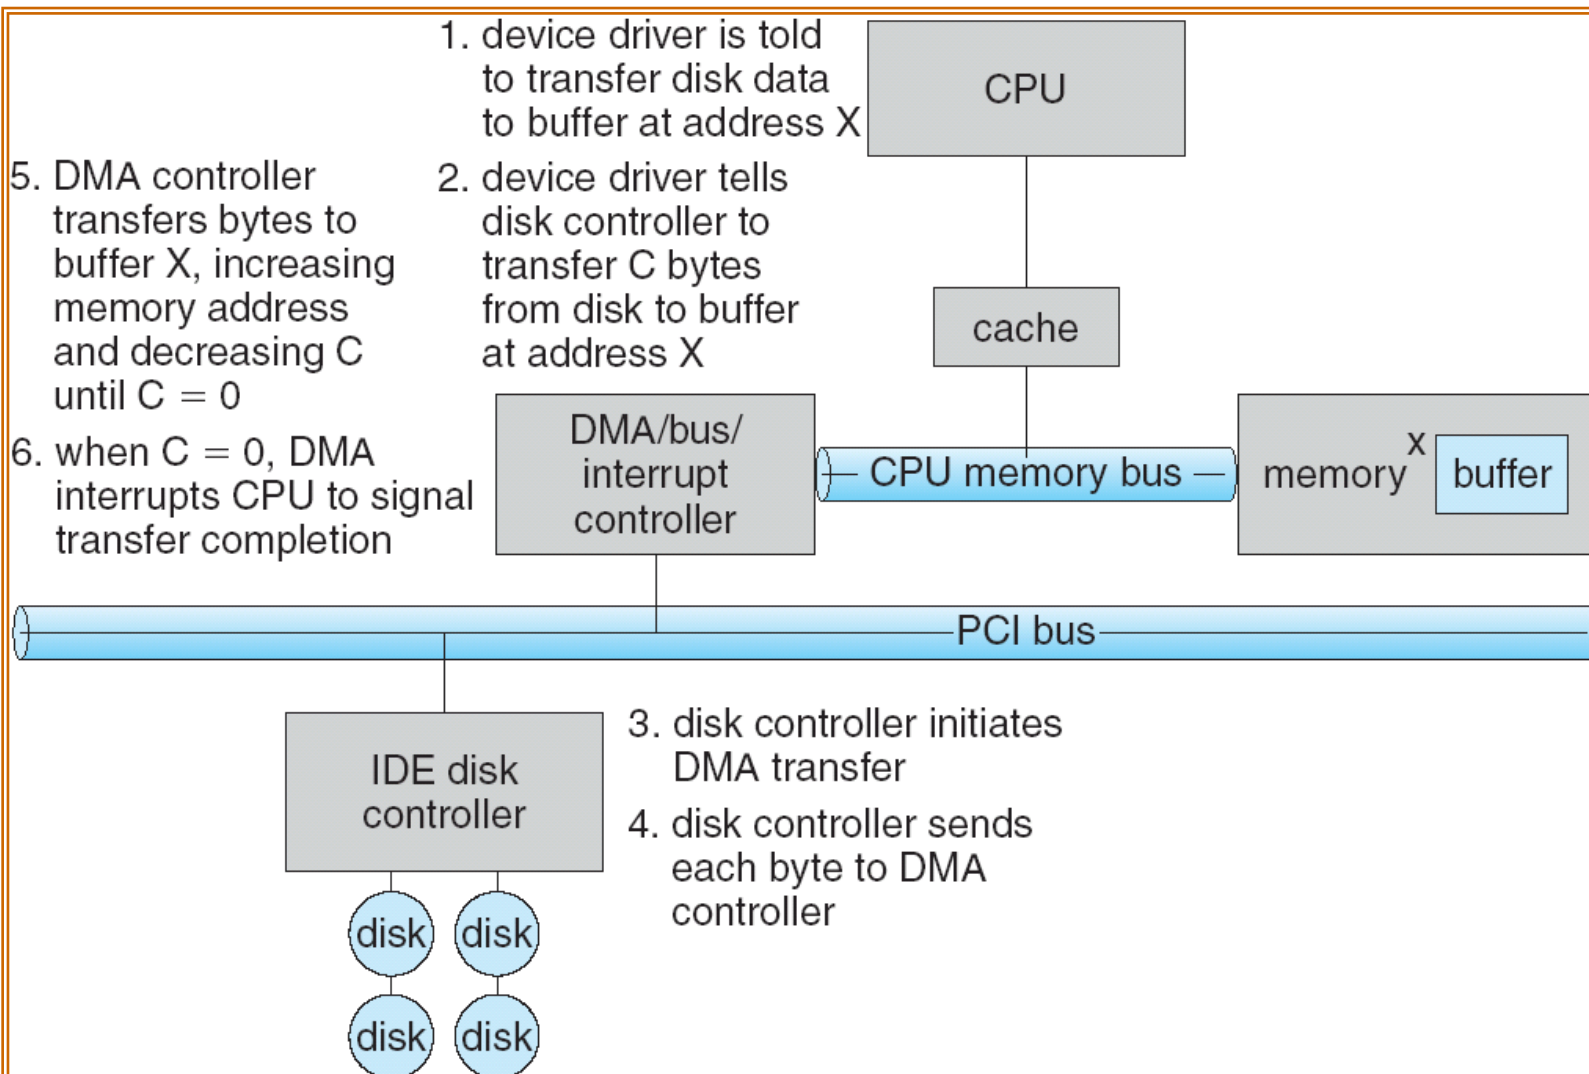


Direct Memory Access

- ❑ Used to avoid **programmed I/O** for large data movement
- ❑ Requires **DMA** controller
- ❑ Bypasses CPU to transfer data directly between I/O device and memory



Six Step Process to Perform DMA Transfer

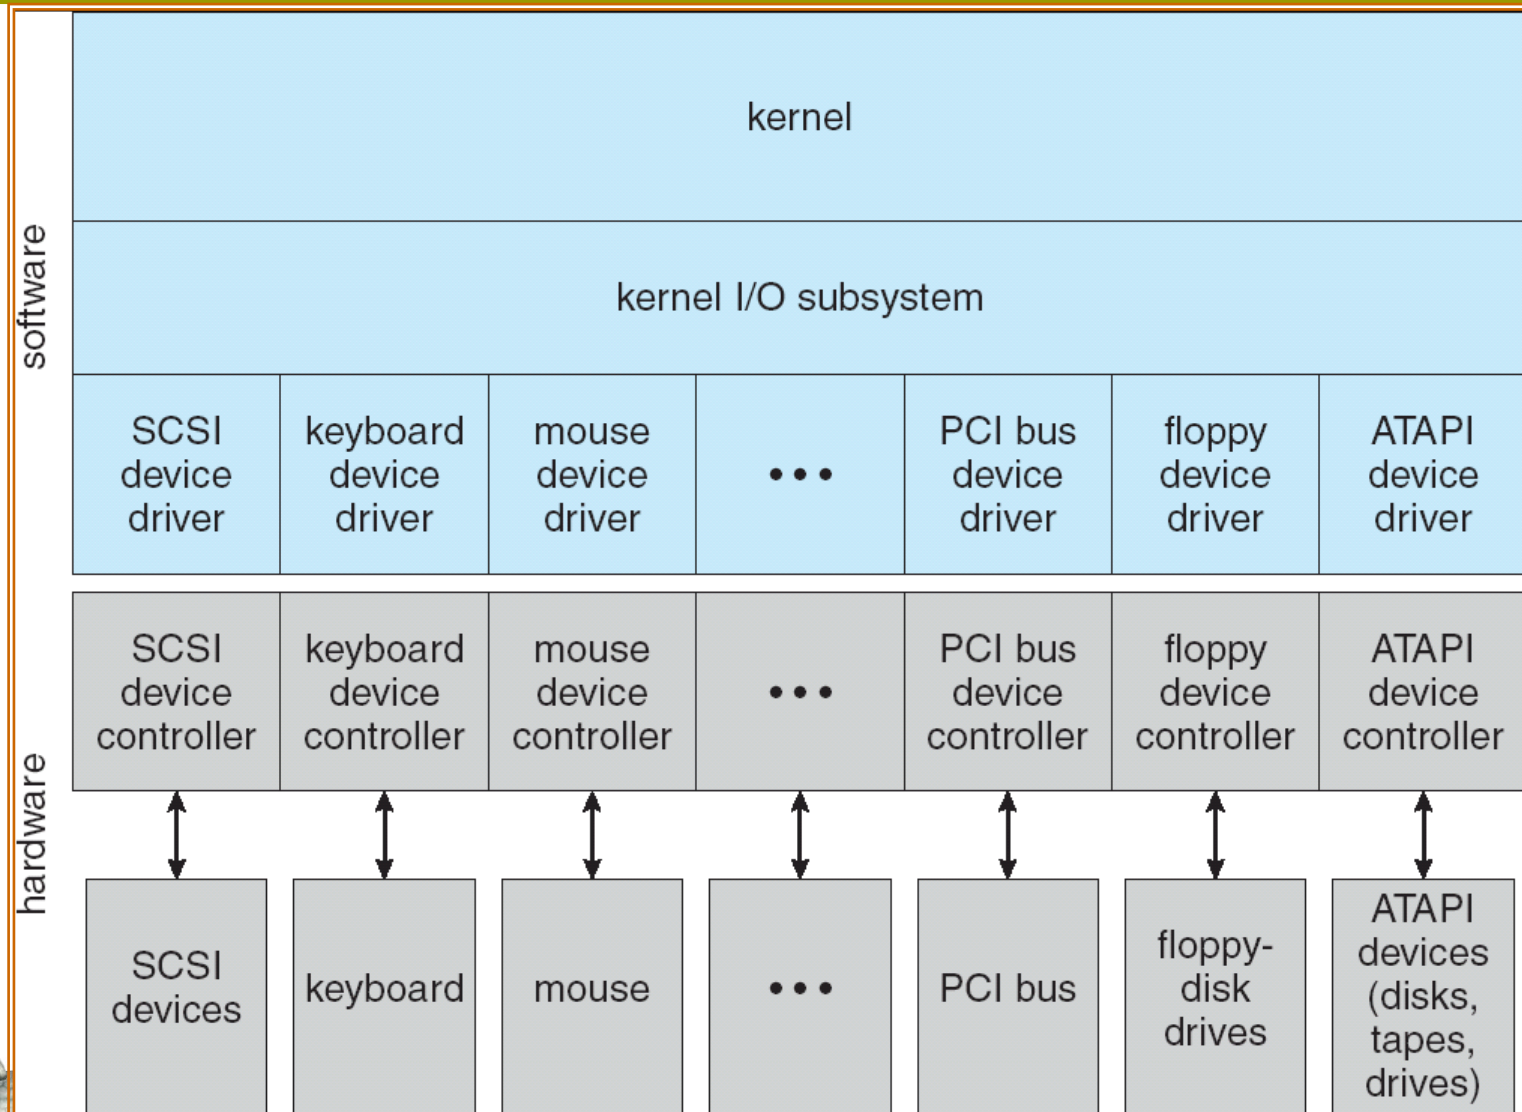


Application I/O Interface

- ❑ I/O system calls **encapsulate** device behaviors in generic classes
- ❑ Device-driver layer **hides differences** among I/O controllers from kernel
- ❑ Devices vary in many dimensions
 - **Character-stream or block**
 - **Sequential or random-access**
 - **Sharable or dedicated**
 - **Speed of operation**
 - **read-write, read only, or write only**



A Kernel I/O Structure



Characteristics of I/O Devices

aspect	variation	example
data-transfer mode	character block	terminal disk
access method	sequential random	modem CD-ROM
transfer schedule	synchronous asynchronous	tape keyboard
sharing	dedicated sharable	tape keyboard
device speed	latency seek time transfer rate delay between operations	
I/O direction	read only write only read–write	CD-ROM graphics controller disk



Block and Character Devices

- ❑ **Block devices** include disk drives
 - Commands include **read**, **write**, **seek**
 - Raw I/O or file-system access
 - Memory-mapped file access possible
- ❑ **Character devices** include keyboards, mice, serial ports
 - Commands include **get**, **put**
 - Libraries layered on top allow line editing



Network Devices

- ❑ Varying enough from block and character to have own interface
- ❑ Unix and Windows NT/9x/2000 include socket interface
 - Separates network protocol from network operation
 - Includes `select` functionality
- ❑ Approaches vary widely (pipes, FIFOs, streams, queues, mailboxes)



Clocks and Timers

- ❑ Provide current time, elapsed time, timer
- ❑ **Programmable interval timer** used for timings, periodic interrupts
- ❑ `ioctl` (on UNIX) covers odd aspects of I/O such as clocks and timers

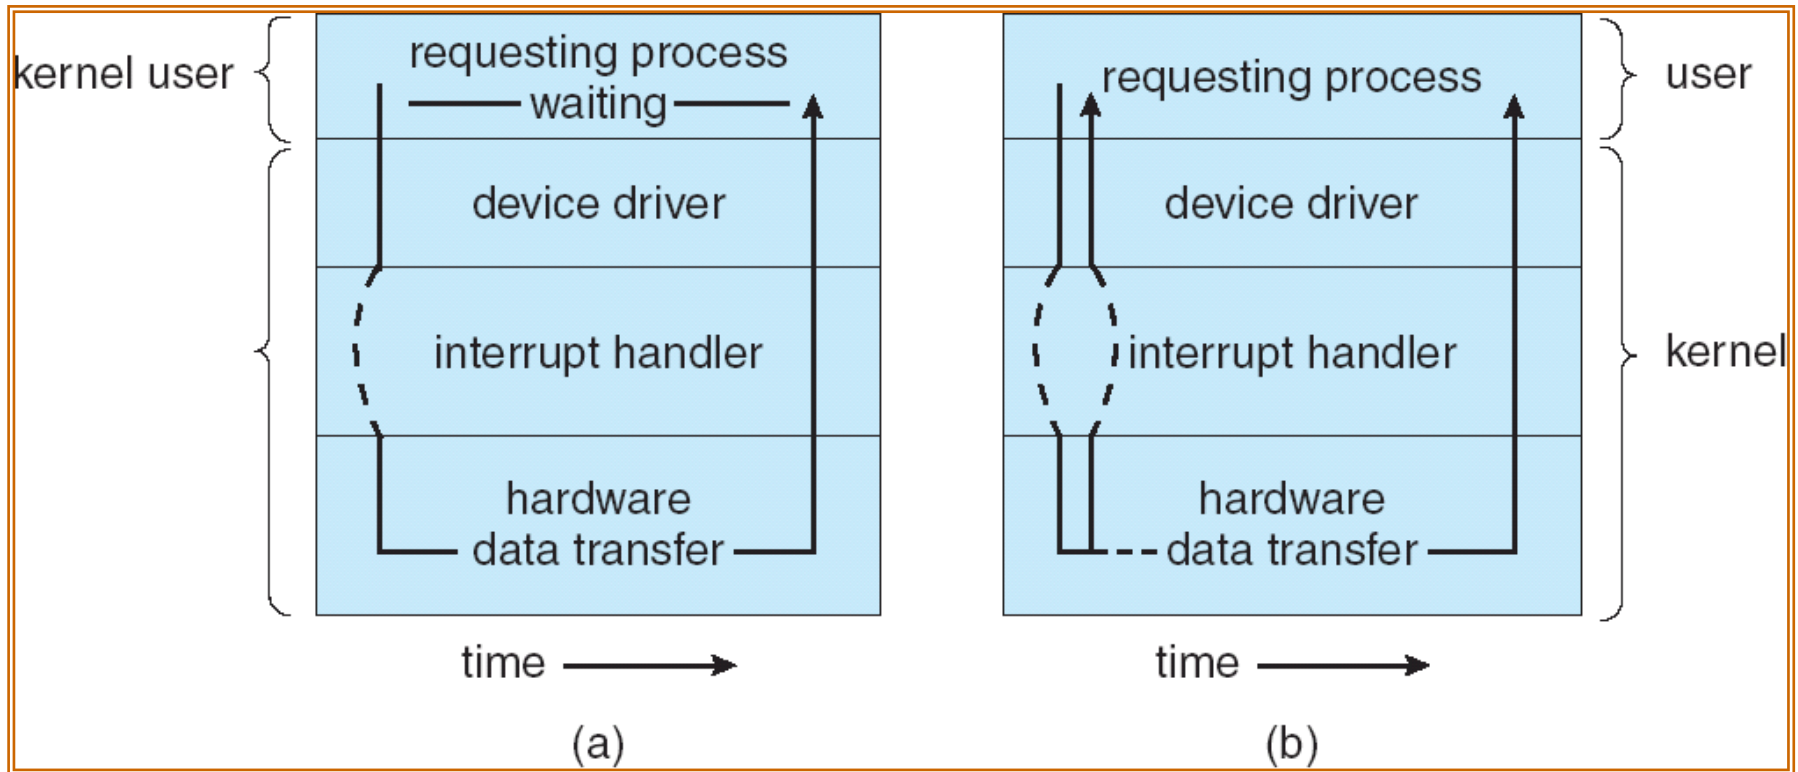


Blocking and Nonblocking I/O

- ❑ **Blocking** - process **suspended** until I/O completed
 - Easy to use and understand
 - Insufficient for some needs
- ❑ **Nonblocking** - I/O call returns as much as available
 - User interface, data copy (buffered I/O)
 - Implemented via multi-threading
 - Returns quickly with **count of bytes read or written**
- ❑ **Asynchronous** - process runs while I/O executes
 - Difficult to use
 - I/O subsystem signals process when I/O completed



Two I/O Methods



Synchronous

Asynchronous

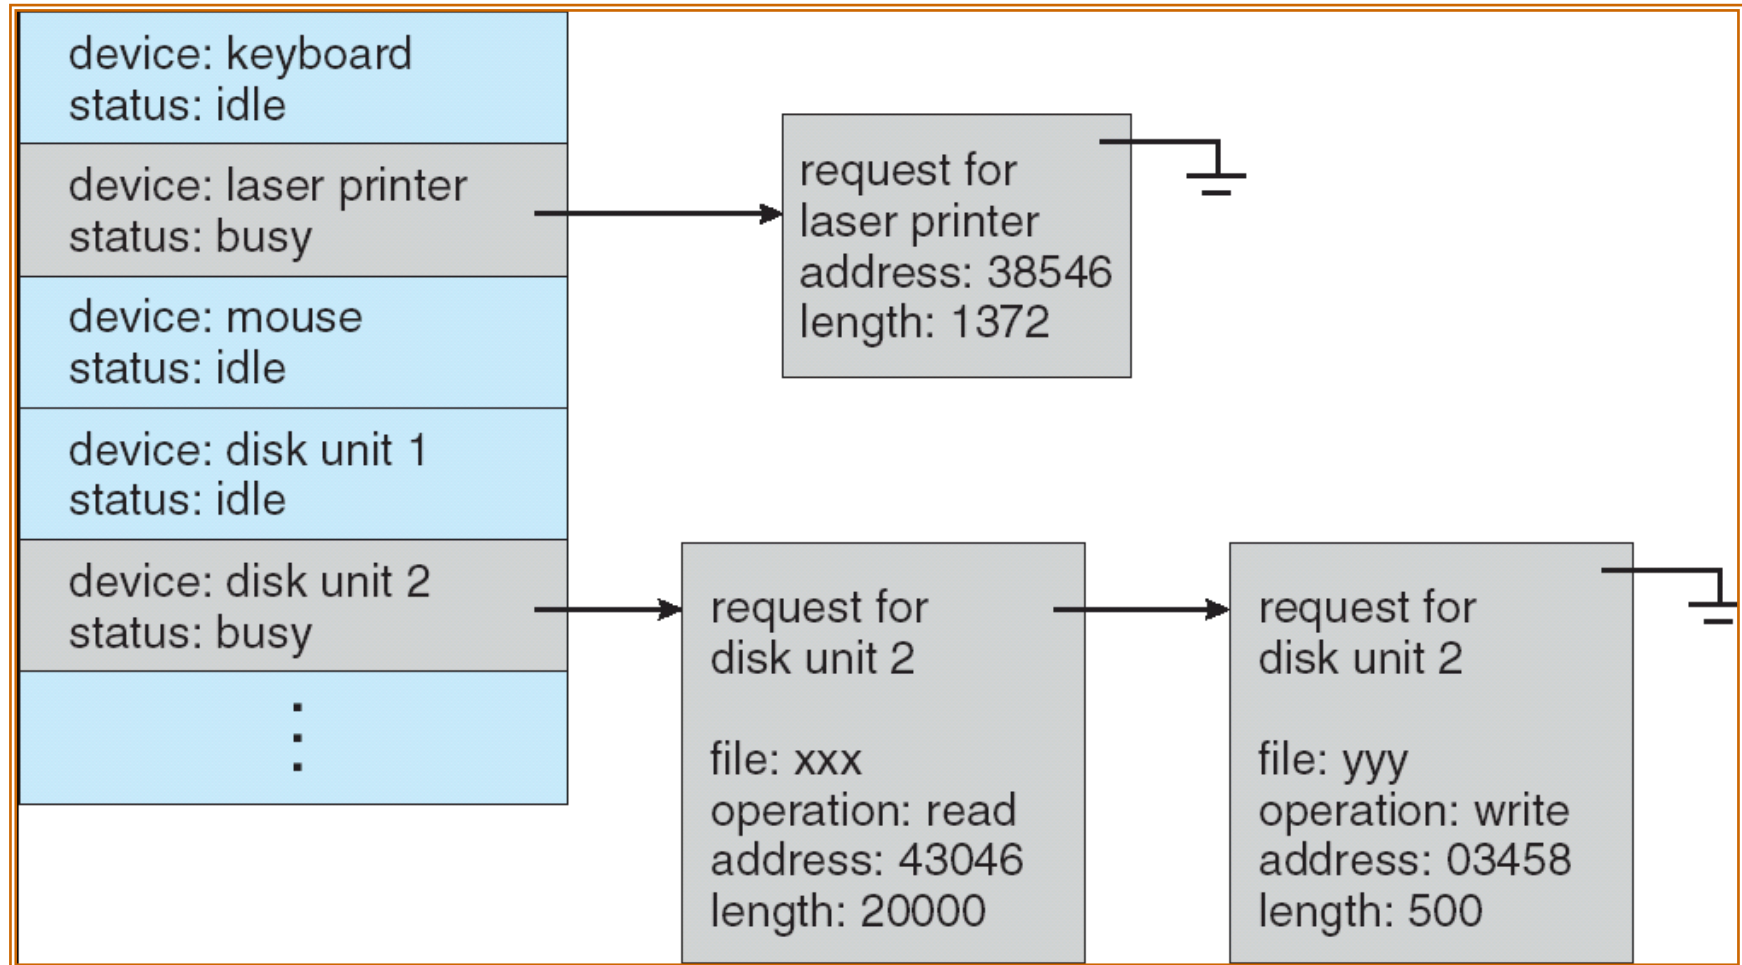


Kernel I/O Subsystem


- ❑ **Scheduling**
 - Some I/O request ordering via per-device queue
 - Some OSs try fairness
- ❑ **Buffering** - store data **in memory** while transferring between devices
 - To cope with device speed **mismatch**
 - To cope with device transfer **size mismatch**
 - **To maintain “copy semantics”**



Device-status Table



Sun Enterprise 6000 Device-Transfer Rates



Kernel I/O Subsystem

- ❑ **Caching** - fast memory holding copy of data
 - Always just a copy (buffer)
 - Key to performance
- ❑ **Spooling** - hold output for a device
 - If device can serve only one request at a time
 - i.e., Printing
- ❑ **Device reservation** - provides exclusive access to a device
 - System calls for allocation and deallocation
 - Watch out for deadlock



Error Handling

- ❑ OS can recover from disk read, device unavailable, transient write failures
- ❑ Most return an **error number** or code when I/O request fails
- ❑ System **error logs** hold problem reports



Improving Performance

- ❑ Reduce number of context switches
- ❑ Reduce data copying
- ❑ Reduce interrupts by using large transfers, smart controllers, polling
- ❑ Use DMA
- ❑ Balance CPU, memory, bus, and I/O performance for highest throughput



End of Chapter 13

